

A Comparison of Language Modeling and Translation as Multilingual Pretraining Objectives

Zihao Li,¹ Shaoxiong Ji,^{*1} Timothee Mickus,^{*1} Vincent Segonne,² and Jörg Tiedemann¹

¹University of Helsinki ²Université Bretagne Sud

firstname.lastname@{¹helsinki.fi, ²univ-ubs.fr}

Abstract

Pretrained language models (PLMs) display impressive performances and have captured the attention of the NLP community. Establishing best practices in pretraining has, therefore, become a major focus of NLP research, especially since insights gained from monolingual English models may not necessarily apply to more complex multilingual models. One significant caveat of the current state of the art is that different works are rarely comparable: they often discuss different parameter counts, training data, and evaluation methodology.

This paper proposes a comparison of multilingual pretraining objectives in a controlled methodological environment. We ensure that training data and model architectures are comparable, and discuss the downstream performances across 6 languages that we observe in probing and fine-tuning scenarios. We make two key observations: (1) the architecture dictates which pretraining objective is optimal; (2) multilingual translation is a very effective pretraining objective under the right conditions. We make our code, data, and model weights available at <https://github.com/Helsinki-NLP/lm-vs-mt>.

1 Introduction

The release of BERT (Devlin et al., 2019) has marked a paradigm shift in the NLP landscape and has ushered in a thorough investment of the NLP research community in developing large language models that can readily be adapted to novel situations. The design, training, and evaluation of these models has become a significant enterprise of its own.

In recent years, that sustained interest has shifted also to encompass multilingual models (e.g., Muenighoff et al., 2022; Alves et al., 2024). There is considerable variation as to how such models are

trained: For instance, some rely on datasets comprising multiple languages without explicit cross-lingual supervision (e.g., Liu et al., 2020), and some use explicit supervision (Xue et al., 2021). One complication that arises from this blossoming field of study is that much of the work being carried out is not directly comparable beyond the raw performances on some well-established benchmark, a procedure which may well be flawed (Gorman and Bedrick, 2019). Avoiding apples-to-oranges comparison requires a methodical approach in strictly comparable circumstances, which is the stance we adopt in this paper.

In short, we focus on two variables—model architecture and pretraining objectives—and set out to train five models in strictly comparable conditions and compare their monolingual performances in three downstream applications: sentiment analysis, named entity recognition, and POS-tagging. The scope of our study spans from encoder-decoder machine translation models, to decoder-only causal language models and encoder-only BERT-like masked language models. We categorize them into **double-stacks (encoder-decoder)** and **single-stacks (encoder-only or decoder-only)** models. We intend to answer two research questions:

- (i) Does the explicit cross-lingual training signal of translation objectives foster better downstream performances in monolingual tasks?
- (ii) Is the optimal choice of architecture independent of the training objective?

There are *a prima facie* reasons to favor either answers to both of these questions. For instance, the success of multilingual pretrained language models (LM) on cross-lingual tasks has been underscored repeatedly (Wu and Dredze, 2019, e.g.), yet explicit alignments such as linear mapping (Wang et al., 2019) and L2 alignment (Cao et al., 2020)

^{*}Equal contribution and corresponding authors.

between source and target languages do not necessarily improve the quality of cross-lingual representations (Wu and Dredze, 2020).

Our experiments provide tentative evidence that insofar as a BART denoising autoencoder architecture is concerned, models pretrained with a translation objective consistently outperform those trained with a denoising objective. However, for single-stack transformers, we observe causal language models to perform well in probing and masked language models to generally outperform translation and causal objectives when fine-tuned on downstream tasks. This leads us to conjecture that the optimal pretraining objective depends on the architecture. Furthermore, the best downstream results we observe appear to stem from a machine-translation system, highlighting that MT encoder-decoder systems might constitute an understudied but potentially very impactful type of pretrained model.

2 Methods and Settings

We start our inquiry by adopting a principled stance: We train strictly comparable models with MT and LM objectives before contrasting their performances on monolingual tasks.

Models and objectives. To allow a systematic evaluation, we train models with various neural network architectures and learning objectives. All models are based on the transformer architecture (Vaswani et al., 2017) and implemented in `fairseq` (Ott et al., 2019). We consider both double-stacks (encoder-decoder) and single-stacks (encoder-only or decoder-only) models.

The two double-stack models are variants of the BART architecture of (Lewis et al., 2020); they are trained either on a straightforward machine translation (MT) objective, using language tokens to distinguish the source, or on the original denoising auto-encoder objective of Lewis et al.. We refer to these two models as **2-LM** and **2-MT** respectively.

We also consider three single-stack models: (i) an encoder-only model trained on the masked language modeling objective (**MLM**) of Devlin et al. (2019); (ii) an autoregressive causal language model (**CLM**), similar to Radford et al. (2019); and (iii) an autoregressive model trained to generate a sentence, followed by its translation in the language specified by a given control token, known as a translation language model (**TLM**) as proposed

by Conneau and Lample (2019).¹ We provide an example datapoint for each pretraining objective in Table 3, Appendix A.

Pretraining conditions. Our core focus is on guaranteeing comparable conditions across the different pretraining objectives we consider. This entails that our datasets need to be doubly structured: both in documents for CLM pretraining; and as aligned bitexts for MT pretraining. Two datasets broadly match these criteria: the UNPC (Ziemski et al., 2016) and OpenSubtitles (OpSub; Tiedemann, 2012) corpora. The choice also narrows down the languages considered in this study: we take the set of languages present in both resources, namely the six languages in UNPC: Arabic (AR), Chinese (ZH), English (EN), French (FR), Russian (RU), and Spanish (ES).

To guarantee that models are trained on the same data, whenever a document is available in multiple languages, we greedily assign it to the least represented language pair thus far and discard all other possible language pairs where it could have contributed; we then discard documents which cannot be used as bitexts. This ensures that all documents are used exactly once for both document-level and bitext-level pretraining objectives. Dataset statistics are shown in Table 4, Appendix B.

To ensure a fair comparison, we control key variables, including tokenization (100k BPE pieces; Sennrich et al., 2016), number of transformer layers (12), hidden dimensions (512), attention heads (8), and feedforward layer dimensions (2048). We perform 600k steps of updates,² using the largest batch size that fits into the GPU memory, deploy distributed training to make a global batch size of 4096, and apply the Adam optimizer (Kingma and Ba, 2017). Owing to the computational requirements, we only train one seed for each of the five types of models considered.

Downstream evaluation. The evaluations encompassed both sequence-level and token-level classification tasks using datasets tailored for sentiment analysis (SA), named entity recognition (NER), part-of-speech (POS) tagging, and natural language inference (NLI).

For SA, we utilized the Amazon review dataset (Hou et al., 2024) in English, Spanish,

¹In this work, we only focus on the causal variant of TLM proposed by Conneau and Lample.

²Improvements in cross-entropy over the validation set were always marginal after this stage.

French, and Chinese. RuReviews (Smetanin and Komarov, 2019) for Russian, and ar_res_reviews (ElSahar and El-Beltagy, 2015) for Arabic. While the datasets for most languages were pre-split, ar_res_reviews required manual division into training, validation, and testing sets, using an 8:1:1 ratio.

For NER, we model the problem as an entity span extraction using a BIO scheme. In practice, we classify tokens into three basic categories: Beginning of an entity (B), Inside an entity (I), or Outside any entity (O). We use the MultiCoNER v2 dataset (Fetahu et al., 2023) for English, Spanish, French, and Chinese, MultiCoNER v1 (Malmasi et al., 2022) for Russian and the AQMAR Wikipedia NER corpus (Mohit et al., 2012a) for Arabic. Simplifying the NER task to these fundamental categories allows us to focus more on assessing the basic entity recognition capabilities of the models without the additional complexity of differentiating numerous entity types, which can vary significantly between languages and datasets.

For POS tagging, we utilized the Universal Dependencies (UD) 2.0 datasets (Nivre et al., 2020), selecting specific corpora tailored to each language to ensure both linguistic diversity and relevance. We select multiple UD treebanks per language, such that each language dataset comprises approximately 160,000 tokens, which are then split into training, validation, and testing segments with an 8:1:1 ratio.

For NLI, we employed the XNLI dataset (Conneau et al., 2018) for the six languages. The XNLI dataset consists of sentence pairs translated from the MultiNLI dataset (Williams et al., 2018) into 15 languages, providing consistent annotations across languages. The task focuses on classifying the relationship between pairs of sentences into one of three categories: Entailment, Contradiction, or Neutral. Unlike the original cross-lingual design of XNLI, we conducted monolingual experiments for each language to evaluate the performance of our models individually in each linguistic context.

Supplementary details regarding data preprocessing for downstream experiments are available in Appendix B.

We evaluate the performances of the encoder output representations for the 2-MT and 2-LM models and of the last hidden representation before the vocabulary projection for the single-stack models.

The evaluation of the models involves two dis-

tinct experimental approaches to test the performance: probing and fine-tuning. In the probing experiments, only the parameters of the classification heads are adjusted. This method primarily tests the raw capability of the pre-trained models’ embeddings to adapt to specific tasks with minimal parameter changes, preserving the underlying pre-trained network structure. Conversely, in the fine-tuning experiments, all parameters of the models are adjusted. This approach allows the entire model to adapt to the specifics of the task, potentially leading to higher performance at the cost of significantly altering the pre-trained weights.

For both experimental approaches, each model is trained for 10 epochs to ensure sufficient learning without overfitting. We optimize parameters with AdamW (Loshchilov and Hutter, 2017), with a constant learning rate of 0.0001 across all tasks and models. This setup was chosen to standardize the training process, providing a fair basis for comparing the performance outcomes across different models and tasks. We reproduce probing and fine-tuning for 5 seeds to ensure stability.

3 Results

Double-stack models. We first compare the performance of 2-LM and 2-MT across several key language processing tasks including SA, NER, POS tagging, and NLI. Results are shown in Tables 1a and 1b. The pretraining objectives play a significant role in shaping the models’ effectiveness. Specifically, 2-MT, which is pretrained with a machine translation objective, consistently outperforms 2-LM, which utilizes a denoising objective. This pattern is consistent across all languages tested after fine-tuning as well as probing.

Single-stack models. Turning to the single-stack models (CLM, MLM, TLM), we find a somewhat more complex picture. In a probing context (cf. Table 2a), we find the CLM to be almost always the most effective, except for NLI in five languages and NER in Arabic, where it performs slightly less favorably compared to the MLM. As for fine-tuning (Table 2b), while the MLM generally ranks first on all POS, NER, and NLI datasets, the TLM is usually effective for SA.³

³However, remark that unlike with the BART-based models, SA results are not stable when we shift metrics from accuracy to F1 (see Tables 6 and 7 in Appendix C). The difference in F1 between the top two models is often ≤ 0.01 , making it difficult to ascertain that one model strictly dominates.

	Setup	Languages					
		EN	ES	FR	ZH	RU	AR
SA	2-LM	42.86±0.84	42.80±0.69	43.00±0.60	40.41±1.02	65.83±0.70	70.88±1.62
	2-MT	46.71 ±0.88	46.64 ±0.55	46.10 ±0.43	43.74 ±0.65	68.79 ±0.42	73.77 ±0.97
NER	2-LM	82.69±0.09	84.74±0.07	82.80±0.06	78.88±0.25	77.93±0.15	85.28±0.22
	2-MT	89.47 ±0.06	90.54 ±0.04	89.41 ±0.10	88.78 ±0.09	83.39 ±0.22	89.70 ±0.18
POS	2-LM	78.85±0.29	78.12±0.25	81.57±0.32	69.09±0.25	77.93±0.12	47.68±0.10
	2-MT	92.22 ±0.14	90.59 ±0.20	95.39 ±0.10	75.87 ±0.17	93.20 ±0.08	61.84 ±0.24
NLI	2-LM	48.56±0.01	49.31±0.01	48.33±0.01	38.81±0.01	48.34±0.01	45.11±0.01
	2-MT	60.50 ±0.01	59.56 ±0.01	59.00 ±0.01	59.01 ±0.01	59.83 ±0.01	59.58 ±0.01

	Setup	Languages					
		EN	ES	FR	ZH	RU	AR
SA	2-LM	52.26±0.55	52.89±0.69	52.99±0.59	48.64±0.36	73.89±0.43	79.74 ±1.36
	2-MT	54.76 ±0.58	55.56 ±0.49	54.75 ±0.42	50.55 ±0.68	74.77 ±0.50	81.49 ±1.49
NER	2-LM	91.13±0.12	91.82±0.21	91.58±0.10	92.30±0.10	85.34±0.39	89.05 ±0.13
	2-MT	93.46 ±0.09	94.22 ±0.09	93.84 ±0.04	93.75 ±0.32	89.07 ±0.11	93.26 ±0.15
POS	2-LM	92.42±0.28	90.41±0.16	95.21±0.13	82.30±0.48	95.36±0.20	69.57 ±0.24
	2-MT	95.98 ±0.08	94.29 ±0.05	98.05 ±0.17	90.18 ±0.15	97.00 ±0.07	74.47 ±0.08
NLI	2-LM	57.76±0.01	57.87±0.01	56.77±0.01	48.05±0.01	56.43±0.01	0.5377±0.01
	2-MT	61.96 ±0.01	61.71 ±0.01	60.09 ±0.01	53.72 ±0.01	59.00 ±0.01	0.5693 ±0.01

(a) Probing

(b) Fine-tuning

Table 1: Accuracy ($\times 100$) of double-stack models (\pm s.d. over 5 runs).

	Setup	Languages					
		EN	ES	FR	ZH	RU	AR
SA	CLM	35.14 ±0.92	35.66 ±1.10	34.14 ±1.63	33.62 ±0.83	57.57 ±1.11	67.71 ±2.24
	MLM	34.26±1.34	34.82±1.58	33.90±1.12	32.52±1.65	54.55±1.86	65.94±3.30
NER	TLM	29.68±2.22	32.20±3.07	32.26±2.34	29.88±4.17	56.45±1.81	64.45±1.81
	CLM	80.27 ±0.12	82.59 ±0.06	80.38 ±0.12	77.92 ±0.28	76.39 ±0.03	84.17±0.08
POS	MLM	78.77±0.02	81.61±0.00	79.11±0.01	70.67±0.10	76.34±0.01	84.29 ±0.13
	TLM	79.10±0.06	81.94±0.13	79.56±0.14	77.26±0.24	76.39 ±0.02	84.26±0.02
NLI	CLM	69.06 ±0.38	70.32 ±0.50	76.67 ±0.44	51.40 ±0.47	59.64 ±0.62	43.49 ±0.40
	MLM	37.92±0.61	44.26±0.11	46.89±0.32	31.16±0.21	34.62±0.16	34.71±0.94
	TLM	62.96±1.02	62.08±1.99	63.89±1.06	50.46±0.53	54.27±0.87	40.94±1.16

(a) Probing

(b) Fine-tuning

Table 2: Accuracy ($\times 100$) of single-stack models (\pm s.d. over 5 runs).

Discussion. A first global observation that we can make for these results is that single-stack and double-stack models appear to behave differently. While the MT objective yields the highest performances for BART-type models, the downstream performances of the TLM do not really stand out compared to the CLM in probing and the MLM in fine-tuning scenarios. It is important to note that the performances stem at least in part from the architecture itself: 2-MT and 2-LM both consistently outperform all single-stack models in probing. However, it is crucial to acknowledge the limitations of our study, as we only conducted one pretraining round for all the objectives. Hence, this evidence should be interpreted as tentative at best.

Fine-tuning also tends to minimize the difference between single-stack and double-stack models—which suggests that the higher quality of double-stack representations could be an artifact of training limitations. Moreover, the relative ranks of the three single-stack models fluctuate much more than what we see for the double-stack models, owing to no little extent to the oftentimes momentous variation across seeds for single-stack models. We therefore conjecture that while a translation objective

can yield a clear training signal towards semantically informed representations, this comes with two caveats: first, the signal can only be leveraged with dedicated separate modeling of source and target (viz. double-stack models); second, this advantage is much less consequential when fine-tuning.

4 Related works

Multilingual foundation models have flourished in recent years (a.o., Conneau and Lample, 2019; Liu et al., 2020; Xue et al., 2021; Kale et al., 2021; Fang et al., 2021; Chi et al., 2021; Alves et al., 2024; Üstün et al., 2024), and with them so have studies of their representations (Conneau et al., 2020; Siddhant et al., 2020; Choudhury and Deshpande, 2021; Fierro and Søgaard, 2022; Hämmel et al., 2023 a.o.). All of these works, however, fail to control for some of the most crucial factors, such as ensuring that all models are trained on comparable amounts of data.

This work is specifically related to Conneau and Lample (2019), which also compares MLM, CLM, and TLM but does not normalize the training data. Another point of comparison is Ji et al. (2024), which studies the impact of MT continued pretrain-

ing in BART on cross-lingual downstream tasks. Monolingual evaluation of multilingual systems has also been broached a.o. by Rust et al. (2021).

5 Conclusion

This paper conducts an empirical study of how pretraining conditions of multilingual models impact downstream performances in probing and fine-tuning scenarios. Despite the inherent limitations that stem from our stringent data requirements, our experiments offer a novel perspective that highlights directions for future inquiry into how multilingual foundation models ought to be pretrained. We observe that double-stack BART-based models fare much better than single-stack models in probing scenarios, but the difference is overall less clear when it comes to fine-tuning. We also find some tentative evidence that translation objectives can be highly effective for model pretraining in precise circumstances: Namely, the most effective model on downstream tasks among those we experimented with is an MT-pretrained BART-like model, which outperforms both a more traditional denoising objective for BART as well as decoder-only CLM and encoder-only MLM models. This would suggest that translation can serve as a powerful pretraining objective, although it is currently under-explored.⁴

Another crucial aspect of our study is that we present strictly comparable models, trained on comparable data, with comparable parameter counts and unified implementations. While this entails some limitations, especially with regard to the scale of models and data used, we nonetheless believe that a strict comparison can help discriminate between the various factors at play in other works. Here, we find clear evidence that CLM pretraining objectives, such as those used in GPT, outperform MLM-based models, such as BERT, in probing scenarios; we are also able to isolate and highlight how the optimal choice of pretraining objective is contingent on the architecture being employed.

For future work, we recommend exploring *multitask learning* during pretraining by combining objectives like translation, denoising, and language modeling; in such cases, models could harness the strengths of each task to become more robust and versatile. Additionally, investigating *training-*

⁴There are reasonable objections against using MT models as pretrained multilingual foundation models—namely, unlike auto-regressive causal language models, their generation capabilities are strictly tied to translation, thereby requiring some degree of multilingualism from end-users.

free evaluation methods can offer insights into a model’s inherent capabilities without the variability introduced by fine-tuning.

Acknowledgments

We thank Alessandro Raganato and our colleagues at the Helsinki-NLP group for useful discussions throughout this project, as well as the three anonymous reviewers for their comments.

This project has received funding from the European Union’s Horizon Europe research and innovation programme under Grant agreement No 101070350 and from UK Research and Innovation (UKRI) under the UK government’s Horizon Europe funding guarantee [grant number 10052546], and partially funded by the French National Research Agency [grant ANR-23-IAS1-0001]. The contents of this publication are the sole responsibility of its authors and do not necessarily reflect the opinion of the European Union.

The authors wish to thank CSC-IT Center for Science, Finland, for the generous computational resources on the Puhti supercomputer and LUMI supercomputer through the LUMI extreme scale access (MOOMIN and LumiNMT). Some of the experiments were performed using the Jean Zay and Adastra clusters from GENCI-IDRIS [grant 2022 A0131013801].

Limitations

This study employs models that are not large in terms of parameters in the era of large language models. Such a constraint potentially hinders the generalizability of our results to much larger architectures that are capable of handling a broader array of linguistic nuances. Furthermore, our study focuses on a small selected group of languages and specific NLP tasks. This focus might limit the applicability of our findings to other linguistic contexts or more complex real-world applications where diverse language phenomena or different task demands play a crucial role.

Another limitation is our reliance on specific corpora. The datasets utilized, while valuable, represent a potential source of selection bias. They may not fully encompass the vast diversity of global language use, thus skewing the model training and evaluation. Such a bias could affect the robustness and effectiveness of the pretrained models when applied to languages that are not well-represented in the training data.

References

- AllSet Learning. 2023. Chinese grammar wiki.
- Duarte M. Alves, José Pombal, Nuno M. Guerreiro, Pedro H. Martins, João Alves, Amin Farajian, Ben Peters, Ricardo Rei, Patrick Fernandes, Sweta Agrawal, Pierre Colombo, José G. C. de Souza, and André F. T. Martins. 2024. Tower: An open multilingual large language model for translation-related tasks.
- Steven Cao, Nikita Kitaev, and Dan Klein. 2020. Multi-lingual alignment of contextual word representations. In *International Conference on Learning Representations*.
- Zewen Chi, Li Dong, Shuming Ma, Shaohan Huang, Saksham Singhal, Xian-Ling Mao, Heyan Huang, Xia Song, and Furu Wei. 2021. mT6: Multilingual pretrained text-to-text transformer with translation pairs. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 1671–1683, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.
- Monojit Choudhury and Amit Deshpande. 2021. How linguistically fair are multilingual pre-trained language models? *Proceedings of the AAAI Conference on Artificial Intelligence*, 35(14):12710–12718.
- Alexis Conneau and Guillaume Lample. 2019. Cross-lingual language model pretraining. In *Advances in Neural Information Processing Systems*, volume 32. Curran Associates, Inc.
- Alexis Conneau, Ruty Rinott, Guillaume Lample, Adina Williams, Samuel Bowman, Holger Schwenk, and Veselin Stoyanov. 2018. XNLI: Evaluating cross-lingual sentence representations. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 2475–2485, Brussels, Belgium. Association for Computational Linguistics.
- Alexis Conneau, Shijie Wu, Haoran Li, Luke Zettlemoyer, and Veselin Stoyanov. 2020. Emerging cross-lingual structure in pretrained language models. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 6022–6034, Online. Association for Computational Linguistics.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota. Association for Computational Linguistics.
- Hady ElSahar and Samhaa R El-Beltagy. 2015. Building large arabic multi-domain resources for sentiment analysis. In *International conference on intelligent text processing and computational linguistics*, pages 23–34. Springer.
- Yuwei Fang, Shuohang Wang, Zhe Gan, Siqi Sun, and Jingjing Liu. 2021. Filter: An enhanced fusion method for cross-lingual language understanding. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pages 12776–12784.
- Besnik Fetahu, Zhiyu Chen, Sudipta Kar, Oleg Rokhlenko, and Shervin Malmasi. 2023. Multi-CoNER v2: a large multilingual dataset for fine-grained and noisy named entity recognition. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 2027–2051, Singapore. Association for Computational Linguistics.
- Constanza Fierro and Anders Søgaard. 2022. Factual consistency of multilingual pretrained language models. In *Findings of the Association for Computational Linguistics: ACL 2022*, pages 3046–3052, Dublin, Ireland. Association for Computational Linguistics.
- Kyle Gorman and Steven Bedrick. 2019. We need to talk about standard splits. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 2786–2791, Florence, Italy. Association for Computational Linguistics.
- Bruno Guillaume, Marie-Catherine de Marneffe, and Guy Perrier. 2019. Conversion et améliorations de corpus du français annotés en Universal Dependencies [conversion and improvement of Universal Dependencies French corpora]. *Traitements Automatiques des Langues*, 60(2):71–95.
- Katharina Hämerl, Alina Fastowski, Jindřich Líbovický, and Alexander Fraser. 2023. Exploring anisotropy and outliers in multilingual language models for cross-lingual semantic sentence similarity. In *Findings of the Association for Computational Linguistics: ACL 2023*, pages 7023–7037, Toronto, Canada. Association for Computational Linguistics.
- Yupeng Hou, Jiacheng Li, Zhankui He, An Yan, Xiush Chen, and Julian McAuley. 2024. Bridging language and items for retrieval and recommendation. *arXiv preprint arXiv:2403.03952*.
- Shaoxiong Ji, Timothee Mickus, Vincent Segonne, and Jörg Tiedemann. 2024. Can machine translation bridge multilingual pretraining and cross-lingual transfer learning? In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, pages 2809–2818, Torino, Italia. ELRA and ICCL.
- Mihir Kale, Aditya Siddhant, Rami Al-Rfou, Linting Xue, Noah Constant, and Melvin Johnson. 2021. nmT5 - is parallel data still relevant for pre-training massively multilingual language models? In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language*

- Processing (Volume 2: Short Papers)*, pages 683–691, Online. Association for Computational Linguistics.
- Diederik P. Kingma and Jimmy Ba. 2017. Adam: A method for stochastic optimization.
- John Lee, Herman Leung, and Keying Li. 2017. Towards Universal Dependencies for learner Chinese. In *Proceedings of the NoDaLiDa 2017 Workshop on Universal Dependencies (UDW 2017)*, pages 67–71, Gothenburg, Sweden. Association for Computational Linguistics.
- Mike Lewis, Yinhan Liu, Naman Goyal, Marjan Ghazvininejad, Abdelrahman Mohamed, Omer Levy, Veselin Stoyanov, and Luke Zettlemoyer. 2020. BART: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 7871–7880, Online. Association for Computational Linguistics.
- Yixuan Li, Gerdes Kim, Guillaume Bruno, and Dan Zeman. 2022. Ud chinese patentchar.
- Yinhan Liu, Jiatao Gu, Naman Goyal, Xian Li, Sergey Edunov, Marjan Ghazvininejad, Mike Lewis, and Luke Zettlemoyer. 2020. Multilingual denoising pre-training for neural machine translation. *Transactions of the Association for Computational Linguistics*, 8:726–742.
- Ilya Loshchilov and Frank Hutter. 2017. Decoupled weight decay regularization. *arXiv preprint arXiv:1711.05101*.
- Olga Lyashevskaya, Olga Rudina, Natalia Vlasova, and Anna Zhuravleva. 2018. Ud russian taiga.
- Shervin Malmasi, Anjie Fang, Besnik Fetahu, Sudipta Kar, and Oleg Rokhlenko. 2022. MultiCoNER: A large-scale multilingual dataset for complex named entity recognition. In *Proceedings of the 29th International Conference on Computational Linguistics*, pages 3798–3809, Gyeongju, Republic of Korea. International Committee on Computational Linguistics.
- Ryan McDonald, Joakim Nivre, Yvonne Quirmbach-Brundage, Yoav Goldberg, Dipanjan Das, Kuzman Ganchev, Keith Hall, Slav Petrov, Hao Zhang, Oscar Täckström, Claudia Bedini, Núria Bertomeu Castelló, and Jungmee Lee. 2013. Universal Dependency annotation for multilingual parsing. In *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 92–97, Sofia, Bulgaria. Association for Computational Linguistics.
- Behrang Mohit, Nathan Schneider, Rishav Bhowmick, Kemal Oflazer, and Noah A. Smith. 2012a. Recall-oriented learning of named entities in Arabic Wikipedia. In *Proceedings of the 13th Conference of the European Chapter of the Association for Computational Linguistics*, pages 162–173, Avignon, France. Association for Computational Linguistics.
- Behrang Mohit, Nathan Schneider, Rishav Bhowmick, Kemal Oflazer, and Noah A. Smith. 2012b. Recall-oriented learning of named entities in arabic wikipedia. In *Proceedings of the 13th Conference of the European Chapter of the Association for Computational Linguistics*, pages 162–173.
- Niklas Muennighoff, Thomas Wang, Lintang Sutawika, Adam Roberts, Stella Biderman, Teven Le Scao, M Saiful Bari, Sheng Shen, Zheng-Xin Yong, Hailey Schoelkopf, et al. 2022. Crosslingual generalization through multitask finetuning. *arXiv preprint arXiv:2211.01786*.
- Joakim Nivre, Željko Agić, Lars Ahrenberg, Lene Antonsen, Maria Jesus Aranzabe, Masayuki Asahara, Luma Ateyah, Mohammed Attia, Aitziber Atutxa, Elena Badmaeva, Miguel Ballesteros, Esha Banerjee, Sebastian Bank, John Bauer, Kepa Benigoetxea, Riyaz Ahmad Bhat, Eckhard Bick, Cristina Bosco, Gosse Bouma, Sam Bowman, Aljoscha Burchardt, Marie Candito, Gauthier Caron, Gülsen Cebiroğlu Eryiğit, Giuseppe G. A. Celano, Savas Cetin, Fabricio Chalub, Jinho Choi, Yongseok Cho, Silvie Cinková, Çağrı Çöltekin, Miriam Connor, Marie-Catherine de Marneffe, Valeria de Paiva, Arantza Diaz de Ilarrazza, Kaja Dobrovoljc, Timothy Dozat, Kira Droganova, Marhaba Eli, Ali Elkahky, Tomaž Erjavec, Richárd Farkas, Hector Fernandez Alcalde, Jennifer Foster, Cláudia Freitas, Katarína Gajdošová, Daniel Galbraith, Marcos Garcia, Filip Ginter, Iakes Goenaga, Koldo Gojenola, Memduh Gökirmak, Yoav Goldberg, Xavier Gómez Guinovart, Berta González Saavedra, Matias Grioni, Normunds Grūžītis, Bruno Guillaume, Nizar Habash, Jan Hajčík, Jan Hajčík jr., Linh Hà Mý, Kim Harris, Dag Haug, Barbora Hladká, Jaroslava Hlaváčová, Petter Hohle, Radu Ion, Elena Irimia, Anders Johannsen, Fredrik Jørgensen, Hüner Kaşkara, Hiroshi Kanayama, Jenna Kanerva, Tolga Kayadelen, Václava Kettnerová, Jesse Kirchner, Natalia Kotysba, Simon Krek, Sookyung Kwak, Veronika Laippala, Lorenzo Lambertino, Tatiana Lando, Phôêng Lê Hông, Alessandro Lenci, Saran Lertpradit, Herman Leung, Cheuk Ying Li, Josie Li, Nikola Ljubešić, Olga Loginova, Olga Lyashevskaya, Teresa Lynn, Vivien Macketanz, Aibek Makazhanov, Michael Mandl, Christopher Manning, Ruli Manurung, Cătălina Mărănduc, David Mareček, Katrin Marheinecke, Héctor Martínez Alonso, André Martins, Jan Mašek, Yuji Matsumoto, Ryan McDonald, Gustavo Mendonça, Anna Missilä, Virginica Mititelu, Yusuke Miyao, Simonetta Montemagni, Amir More, Laura Moreno Romero, Shunsuke Mori, Bohdan Moskalevskyi, Kadri Muischnek, Nina Mustafina, Kaili Müürisepp, Pinkey Nainwani, Anna Nedoluzhko, Lôêng Nguyẽn Thị, Huyễn Nguyẽn Thị Minh, Vitaly Nikolaev, Rattima Nitisoroj, Hanna Nurmi, Stina Ojala, Petya Osenova, Lilja Øvreliid, Elena Pascual, Marco Passarotti, Cenel-Augusto Perez, Guy Perrier, Slav Petrov, Jussi Piitulainen, Emily Pitler, Barbara Plank, Martin Popel, Lauma Pretkalniņa, Prokopis Prokopidis, Tiina Puolakainen, Sampo Pyysalo, Alexandre Rademaker,

- Livy Real, Siva Reddy, Georg Rehm, Larissa Rinaldi, Laura Rituma, Rudolf Rosa, Davide Rovati, Shadi Saleh, Manuela Sanguinetti, Baiba Saulīte, Yanin Sawanakunanon, Sebastian Schuster, Djamel Seddah, Wolfgang Seeker, Mojgan Seraji, Lena Shakurova, Mo Shen, Atsuko Shimada, Muh Shohibussirri, Natalia Silveira, Maria Simi, Radu Simionescu, Katalin Simkó, Mária Šimková, Kiril Simov, Aaron Smith, Antonio Stella, Jana Strnadová, Alane Suhr, Umut Sulubacak, Zsolt Szántó, Dima Taji, Takaaki Tanaka, Trond Trosterud, Anna Trukhina, Reut Tsarfaty, Francis Tyers, Sumire Uematsu, Zdeňka Urešová, Larraitz Uria, Hans Uszkoreit, Gertjan van Noord, Viktor Varga, Veronika Vincze, Jonathan North Washington, Zhuoran Yu, Zdeněk Žabokrtský, Daniel Zeman, and Hanzhi Zhu. 2017. Universal dependencies 2.0 – CoNLL 2017 shared task development and test data. LINDAT/CLARIAH-CZ digital library at the Institute of Formal and Applied Linguistics (ÚFAL), Faculty of Mathematics and Physics, Charles University.
- Joakim Nivre, Marie-Catherine De Marneffe, Filip Ginter, Jan Hajčík, Christopher D Manning, Sampo Pyysalo, Sebastian Schuster, Francis Tyers, and Daniel Zeman. 2020. Universal dependencies v2: An evergrowing multilingual treebank collection. *arXiv preprint arXiv:2004.10643*.
- Myle Ott, Sergey Edunov, Alexei Baevski, Angela Fan, Sam Gross, Nathan Ng, David Grangier, and Michael Auli. 2019. fairseq: A fast, extensible toolkit for sequence modeling. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics (Demonstrations)*, pages 48–53, Minneapolis, Minnesota. Association for Computational Linguistics.
- Peng Qi, Koichi Yasuoka, and Dan Zeman. 2019. Ud-chinese gsdsimp.
- Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, Ilya Sutskever, et al. 2019. Language models are unsupervised multitask learners. *OpenAI blog*, 1(8):9.
- Phillip Rust, Jonas Pfeiffer, Ivan Vulić, Sebastian Ruder, and Iryna Gurevych. 2021. How good is your tokenizer? on the monolingual performance of multilingual language models. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 3118–3135, Online. Association for Computational Linguistics.
- Rico Sennrich, Barry Haddow, and Alexandra Birch. 2016. Neural machine translation of rare words with subword units. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1715–1725, Berlin, Germany. Association for Computational Linguistics.
- Aditya Siddhant, Melvin Johnson, Henry Tsai, Naveen Ari, Jason Riesa, Ankur Bapna, Orhan Firat, and Karthik Raman. 2020. Evaluating the cross-lingual effectiveness of massively multilingual neural machine translation. In *Proceedings of the AAAI conference on artificial intelligence*, volume 34, pages 8854–8861.
- Sergey Smetanin and Michail Komarov. 2019. Sentiment analysis of product reviews in russian using convolutional neural networks. In *2019 IEEE 21st Conference on Business Informatics (CBI)*, volume 01, pages 482–486.
- Jörg Tiedemann. 2012. Parallel data, tools and interfaces in opus. In *Proceedings of LREC*, volume 2012, pages 2214–2218.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. In *Advances in Neural Information Processing Systems*, volume 30. Curran Associates, Inc.
- Yuxuan Wang, Wanxiang Che, Jiang Guo, Yijia Liu, and Ting Liu. 2019. Cross-lingual bert transformation for zero-shot dependency parsing. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 5721–5727.
- Adina Williams, Nikita Nangia, and Samuel Bowman. 2018. A broad-coverage challenge corpus for sentence understanding through inference. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 1112–1122, New Orleans, Louisiana. Association for Computational Linguistics.
- Tak-sum Wong, Kim Gerdes, Herman Leung, and John Lee. 2017. Quantitative comparative syntax on the Cantonese-Mandarin parallel dependency treebank. In *Proceedings of the Fourth International Conference on Dependency Linguistics (Depling 2017)*, pages 266–275, Pisa, Italy. Linköping University Electronic Press.
- Shijie Wu and Mark Dredze. 2019. Beto, bentz, becas: The surprising cross-lingual effectiveness of bert. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 833–844.
- Shijie Wu and Mark Dredze. 2020. Do explicit alignments robustly improve multilingual encoders? In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 4471–4482.
- Linting Xue, Noah Constant, Adam Roberts, Mihir Kale, Rami Al-Rfou, Aditya Siddhant, Aditya Barua, and Colin Raffel. 2021. mT5: A massively multilingual pre-trained text-to-text transformer. In *Proceedings*

of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, pages 483–498, Online. Association for Computational Linguistics.

Amir Zeldes. 2017. The GUM corpus: Creating multilayer resources in the classroom. *Language Resources and Evaluation*, 51(3):581–612.

Dan Zeman, Kirian Guiller, and Bruno Guillaume. 2023. Ud chinese beginner.

Otakar Smrž, Viktor Bielický, Iveta Kouřilová, Jakub Kráčmar, Zemánek. 2008. Dependency treebank : A word on the million words.

Michał Ziemska, Marcin Junczys-Dowmunt, and Bruno Pouliquen. 2016. The United Nations parallel corpus v1.0. In *Proceedings of the Tenth International Conference on Language Resources and Evaluation (LREC'16)*, pages 3530–3534, Portorož, Slovenia. European Language Resources Association (ELRA).

Ahmet Üstün, Viraat Aryabumi, Zheng-Xin Yong, Wei-Yin Ko, Daniel D'souza, Gbemileke Onilude, Neel Bhandari, Shivalika Singh, Hui-Lee Ooi, Amr Kayid, Freddie Vargas, Phil Blunsom, Shayne Longpre, Niklas Muennighoff, Marzieh Fadaee, Julia Kreutzer, and Sara Hooker. 2024. Aya model: An instruction finetuned open-access multilingual language model. *arXiv preprint arXiv:2402.07827*.

A Overview of pretraining objectives

Table 3 displays an example data point for all pre-training objectives we consider. In principle, the CLM is a document-level objective, i.e., the full document would be used as an input rather than the two sentences we show here.

B Datasets statistics

An overview of the volume of data available for pretraining is displayed in Table 4. The majority of the data were used for training.

In Table 5, we present an overview of the datasets used for downstream evaluation.

C Detailed results

In Table 6 and Table 7, we present the macro-f1 score of models in the downstream evaluation.

Objective	Source input	Target output
2-LM	_D'autres _mesures _de _ce _type _vont _être [MASK] [MASK], _en _coopération _avec _d'autres _associations _de _Rom s, _de _Sin tis _et _de [MASK] _du _voyage _(`` C am min anti ''). </s>	<s> _D'autres _mesures _de _ce _type _vont _être _appliquées, _en _coopération _avec _d'autres _associations _de _Rom s, _de _Sin tis _et _de _gens _du _voyage _(`` C am min anti ''). </s>
2-MT	<fr> _D'autres _mesures _de _ce _type _vont _être _appliquées, _en _coopération _avec _d'autres _associations _de _Rom s, _de _Sin tis _et _de _gens _du _voyage _(`` C am min anti '').	<s> _Other_similar_measures_are_going_to_be_taken_in_cooperation_with_other_Roma,_Sin_ti_and_Travel_lers_(`` C am min anti '')_associations.</s>
CLM	... _Divers _accords _ad _hoc _ont _été _conclus _à _cet _effet _par _le _Ministère _de _l'éducation _et _l'association _Op era _Nom ad i. _D'autres _mesures _de _ce _type _vont _être _appliquées, _en _coopération _avec _d'autres _associations _de _Rom s, _de _Sin tis _et _de _gens _du _voyage _(`` C am min anti ''). _accords _ad _hoc _ont _été _conclus _à _cet _effet _par _le _Ministère _de _l'éducation _et _l'association _Op era _Nom ad i. _D'autres _mesures _de _ce _type _vont _être _appliquées, _en _coopération _avec _d'autres _associations _de _Rom s, _de _Sin tis _et _de _gens _du _voyage _(`` C am min anti ''). ...
TLM	_D'autres _mesures _de _ce _type _vont _être _appliquées, _en _coopération _avec _d'autres _associations _de _Rom s, _de _Sin tis _et _de _gens _du _voyage _(`` C am min anti ''). <fr2en> _Other_similar_measures_are_going_to_be_taken_in_cooperation_with_other_Roma,_Sin_ti_and_Travel_lers_(`` C am min anti '')_associations.	_mesures _de _ce _type _vont _être _appliquées, _en _coopération _avec _d'autres _associations _de _Rom s, _de _Sin tis _et _de _gens _du _voyage _(`` C am min anti ''). <fr2en> _Other_similar_measures_are_going_to_be_taken_in_cooperation_with_other_Roma,_Sin_ti_and_Travel_lers_(`` C am min anti '')_associations. </s>
MLM	<s> _D'autres _mesures _de _ce _type _vont _être [MASK] [MASK], _en _coopération _avec _d'autres _associations _de _Rom s, _de _Sin tis _et _de [MASK] _du _voyage _(`` C am min anti ''). </s>	<s> _D'autres _mesures _de _ce _type _vont _être _appliquées, _en _coopération _avec _d'autres _associations _de _Rom s, _de _Sin tis _et _de _gens _du _voyage _(`` C am min anti ''). </s>

Table 3: Overview of the different objectives considered in this study. Top two rows: two-stacks (encoder-decoder) models; bottom three rows: single-stack (encoder-only or decoder-only) models.

	Train	Validation	Test	Total
UNPC	114 376 177	76 303	40 712	114 493 192
OpSub	81 622 353	359 035	77 342	82 058 730
Total	195 998 530	435 338	118 054	196 551 922

Table 4: Number of sentences in pretraining corpora.

Task	Language	Dataset	Class Count	Train	Validation	Test	Total
SA	EN		5	200000	5000	5000	210000
	ES	Amazon Review (Hou et al., 2024)	5	200000	5000	5000	210000
	FR		5	200000	5000	5000	210000
	ZH		5	200000	5000	5000	210000
	RU	RuReviews (Smetanin and Komarov, 2019)	3	85601	2143	2137	89881
NER	AR	ar_res_reviews (ElSahar and El-Beltagy, 2015)	2	6680	835	835	8350
	EN	MultiCoNER v2 (Fetahu et al., 2023)	3	253011	13323	3773671	4040005
	ES	MultiCoNER v2	3	262814	13462	3925900	4202176
	FR	MultiCoNER v2	3	247743	13062	3742924	4003729
	ZH	MultiCoNER v2	3	245606	12816	489605	748027
	RU	MultiCoNER v1 (Malmasi et al., 2022)	3	242384	12787	2061318	2316489
POS	AR	AQMAR Wikipedia NER corpus (Mohit et al., 2012b)	3	57053	8615	8185	73853
	EN	UD_English-GUM (Zeldes, 2017)	16	128391	16070	15554	160015
	ES	UD_Spanish-GSD (McDonald et al., 2013)	16	127459	16916	15645	160020
	FR	UD_French-GSD (Guillaume et al., 2019)	15	127638	16207	16167	160012
	UD_Chinese-Beginner (Zeman et al., 2023; AllSet Learning, 2023)+						
	ZH	UD_Chinese-PUD (Nivre et al., 2017)+	16	128935	15680	15758	160373
	UD_Chinese-HK (Wong et al., 2017)+						
	UD_Chinese-CFL (Lee et al., 2017)+						
	UD_Chinese-PatentChar (Li et al., 2022)+						
	UD_Chinese-GSDSmp (Qi et al., 2019)						
NLI	RU	UD_Russian-Taiga (Lyashevskaya et al., 2018)	16	127647	16175	16184	160006
	AR	UD_Arabic-PADT (Zemánek, 2008)	16	127552	16608	15848	160008
	EN		3	392702	2490	5010	400202
	ES		3	392702	2490	5010	400202
	FR	XNLI (Conneau et al., 2018)	3	392702	2490	5010	400202
	ZH		3	392702	2490	5010	400202
	RU		3	392702	2490	5010	400202
	AR		3	392702	2490	5010	400202

Table 5: Statistics of datasets used for downstream evaluation tasks.

Task	Model	Languages					
		EN	ES	FR	ZH	RU	AR
SA	2-LM	0.4130±0.0118	0.4120±0.0160	0.4166±0.0076	0.3859±0.0156	0.6599±0.0101	0.6343±0.0232
	2-MT	0.4588 ±0.0092	0.4554 ±0.0053	0.4448 ±0.0158	0.4260 ±0.0070	0.6935 ±0.0052	0.6864 ±0.0105
	CLM	0.3183±0.0099	0.3351±0.0198	0.3066±0.0192	0.3104±0.0135	0.5693±0.0107	0.5886±0.0106
	MLM	0.3236±0.0270	0.3188±0.0188	0.3153±0.0088	0.2936±0.0107	0.5434±0.0236	0.5804±0.0104
	TLM	0.2593±0.0298	0.2768±0.0589	0.2528±0.0487	0.2344±0.0539	0.5537±0.0307	0.5487±0.0190
NER	2-LM	0.5830±0.0057	0.5616±0.0070	0.5627±0.0039	0.5653±0.0164	0.4178±0.0100	0.4310±0.0179
	2-MT	0.7778 ±0.0014	0.7660 ±0.0014	0.7716 ±0.0031	0.7871 ±0.0043	0.6551 ±0.0088	0.7311 ±0.0099
	CLM	0.4516±0.0110	0.4213±0.0075	0.4306±0.0131	0.5086±0.0053	0.3004±0.0034	0.3223±0.0054
	MLM	0.3003±0.0017	0.2997±0.0001	0.3021±0.0019	0.3341±0.0108	0.2891±0.0001	0.3094±0.0000
	TLM	0.3485±0.0074	0.3471±0.0152	0.3499±0.0173	0.4876±0.0230	0.2941±0.0015	0.3094±0.0001
POS	2-LM	0.7241±0.0040	0.6607±0.0042	0.6848±0.0074	0.5964±0.0072	0.7427±0.0030	0.4678±0.0016
	2-MT	0.8520 ±0.0065	0.7685 ±0.0203	0.8300 ±0.0017	0.7002 ±0.0029	0.8587 ±0.0055	0.6575 ±0.0032
	CLM	0.5621±0.0069	0.5422±0.0066	0.5568±0.0064	0.3761±0.0148	0.4975±0.0140	0.3040±0.0106
	MLM	0.2157±0.0063	0.1499±0.0055	0.1722±0.0084	0.0717±0.0040	0.1275±0.0080	0.1511±0.0127
	TLM	0.4741±0.0147	0.3759±0.0378	0.3744±0.0153	0.3314±0.0112	0.3798±0.0097	0.2299±0.0215
NLI	2-LM	0.4825±0.0075	0.4901±0.0046	0.4779±0.0102	0.3805±0.0089	0.4804±0.0059	0.4445±0.0126
	2-MT	0.6017 ±0.0105	0.5938 ±0.0119	0.5860 ±0.0087	0.5881 ±0.0031	0.5982 ±0.0025	0.5943 ±0.0053
	CLM	0.3946±0.0479	0.4134±0.0227	0.4068±0.0373	0.3744±0.0400	0.3593±0.0519	0.3978±0.0314
	MLM	0.4464±0.0328	0.4330±0.0145	0.4157±0.0347	0.4208±0.0110	0.4162±0.0251	0.4281±0.0126
	TLM	0.3063±0.0361	0.3573±0.0327	0.3940±0.0240	0.3122±0.0876	0.3892±0.0390	0.3360±0.0477

Table 6: Macro F1 score using probing technique.

Task	Model	Languages					
		EN	ES	FR	ZH	RU	AR
SA	2-LM	0.5213±0.0068	0.5254±0.0083	0.5244±0.0135	0.4739±0.0096	0.7421±0.0059	0.7522±0.0151
	2-MT	0.5407±0.0086	0.5510±0.0084	0.5398±0.0054	0.4956±0.0093	0.7522±0.0056	0.7767±0.0156
	CLM	0.5443±0.0072	0.4446±0.2115	0.5421±0.0089	0.5015±0.0187	0.7553±0.0015	0.5283±0.2328
	MLM	0.5441±0.0107	0.5466±0.0314	0.5348±0.0237	0.4972±0.0142	0.7509±0.0135	0.5695±0.1427
	TLM	0.5358±0.0186	0.5501±0.0128	0.5474±0.0137	0.5069±0.0119	0.7586±0.0057	0.4599±0.0943
NER	2-LM	0.8200±0.0042	0.8092±0.0053	0.8259±0.0035	0.8626±0.0022	0.7215±0.0122	0.7274±0.0093
	2-MT	0.8670±0.0017	0.8651±0.0022	0.8727±0.0018	0.8897±0.0042	0.7934±0.0039	0.8685±0.0046
	CLM	0.7950±0.0064	0.8053±0.0028	0.8099±0.0044	0.8129±0.0021	0.6622±0.0182	0.5994±0.1880
	MLM	0.8635±0.0123	0.8580±0.0142	0.8706±0.0055	0.8739±0.0199	0.7629±0.0172	0.4113±0.2254
	TLM	0.7908±0.0028	0.8024±0.0081	0.8067±0.0047	0.8120±0.0032	0.6758±0.0312	0.3094±0.0000
POS	2-LM	0.8925±0.0039	0.7365±0.0025	0.8496±0.0034	0.8088±0.0059	0.8984±0.0055	0.7769±0.0102
	2-MT	0.9314±0.0024	0.7826±0.0235	0.8866±0.0074	0.8842±0.0059	0.9285±0.0029	0.8660±0.0088
	CLM	0.8752±0.0042	0.7854±0.0024	0.8573±0.0041	0.7906±0.0195	0.8264±0.0104	0.5932±0.0194
	MLM	0.9177±0.0068	0.8079±0.0259	0.8851±0.0019	0.8313±0.0079	0.9226±0.0048	0.8602±0.0132
	TLM	0.8782±0.0045	0.7830±0.0067	0.7421±0.2503	0.7876±0.0271	0.8247±0.0088	0.6201±0.0071
NLI	2-LM	0.5771±0.0067	0.5760±0.0088	0.5658±0.0085	0.4766±0.0058	0.5629±0.0052	0.5350±0.0070
	2-MT	0.6183±0.0054	0.6151±0.0082	0.5991±0.0073	0.5302±0.0086	0.5887±0.0041	0.5678±0.0032
	CLM	0.4240±0.2315	0.5589±0.0355	0.5493±0.0404	0.4729±0.1123	0.5507±0.0265	0.4554±0.1199
	MLM	0.5927±0.0189	0.5719±0.0487	0.5282±0.0964	0.4618±0.0453	0.5775±0.0069	0.5247±0.0221
	TLM	0.4428±0.1751	0.4728±0.1731	0.5345±0.1076	0.4558±0.0722	0.5061±0.0771	0.3816±0.1562

Table 7: Macro F1 score after model fine-tuning.