

MiTTeNS: A Dataset for Evaluating Gender Mistranslation

Kevin Robinson¹ Sneha Kudugunta² Romina Stella¹ Sunipa Dev¹ Jasmijn Bastings¹
Google DeepMind¹ Google Research¹ University of Washington
{kevinrobinson,snehakudugunta,romistella,sunipadev,bastings}@google.com

Abstract

Translation systems, including foundation models capable of translation, can produce errors that result in gender mistranslations, and such errors create potential for harm. To measure the extent of such potential harms when translating into and out of English, we introduce a dataset, MiTTeNS¹, covering 26 languages from a variety of language families and scripts, including several traditionally underrepresented in digital resources. The dataset is constructed with handcrafted passages that target known failure patterns, longer synthetically generated passages, and natural passages sourced from multiple domains. We demonstrate the usefulness of the dataset by evaluating both neural machine translation systems and foundation models, and show that all systems exhibit gender mistranslation and potential harm, even in high resource languages.

1 Introduction

It is well documented that dedicated machine translation systems show forms of gender bias (see Savoldi et al., 2021, for an overview). Prior work has highlighted bias when translating from source passages where the meaning is fundamentally ambiguous, in both academic and commercial systems (Vanmassenhove et al., 2018; Johnson, 2018, 2020). Forms of bias have been demonstrated with carefully constructed unambiguous English passages (Stanovsky et al., 2019), and with linguistic constructions targeting specific language pairs (Cho et al., 2019; Bentivogli et al., 2020; Alhafni et al., 2022; Singh, 2023a,b; Stella, 2021, i.a.).

Recent advances have enabled general-purpose foundation models with powerful multilingual capabilities including translation (Ouyang et al., 2022; OpenAI et al., 2023; Chung et al., 2022; Gemini Team Google, 2023). These models can be used as building blocks in a wide range of products and applications, highlighting the importance of other work on gender bias in natural language processing more broadly (Sun et al., 2019; Costa-jussa, 2019; Stanczak and Augenstein, 2021, i.a.).

Evaluating foundation models raises new challenges of measurement validity, given the wide range of use and potential harms (Weidinger et al., 2023; Shelby et al., 2023). Skew in training data and measures of bias in underlying models may not be reliable predictors or measurements of potential harm in downstream usage (Goldfarb-Tarrant et al., 2021;

Blodgett et al., 2020, 2021). There also remain challenges in empirically measuring performance as systems rapidly improve (Jun, 2023; Krawczyk, 2023), ensuring high quality of service as multilingual capabilities expand (Akter et al., 2023; Yong et al., 2023) and measuring unintentional harms in new system designs (Renduchintala et al., 2021; Costa-jussa et al., 2023).

In this work, we focus on measuring gender mistranslation in both dedicated translation systems and foundation models that can perform translation. Figure 1 illustrates gender mistranslation, and examples of translations that refer to a person in a way that does not reflect the gender identity encoded in the source passage. We focus specifically on gender mistranslation over other harms (Costajussa et al., 2023), and on expanding coverage of language families and scripts at different levels of digital representation (Stanovsky et al., 2019).

IMAGE NOT PROVIDED

Figure 1: Dataset examples targeting passages where gender mistranslation may occur and cause harm. Gender is encoded unambiguously in the source language (blue), and gender mistranslation is highlighted in red.

Adapting evaluation methods to measure gender mistranslation for foundation models presents a few challenges. First, language models are often trained on public internet datasets (Yang et al., 2023; Anil et al., 2023) which can cause contamination and render evaluation sets mined from public data sources ineffective (Kiela et al., 2021). Second, gender is encoded in different ways across languages, making it challenging to scale automated evaluation methods. Automated methods enable faster modeling iteration, but methods commonly used in translation evaluations (eg, BLEU, BLEURT) may fail to capture specific dimensions of harm from gender mistranslation. Finally, the evolving and contested nature of sociocultural norms related to gender make general purpose benchmark methods challenging to develop, particularly for expressions of non-binary gender across linguistic and cultural contexts globally (Dev et al., 2021; Lauscher et al., 2023; Hossain et al., 2023; Cao and Daume III, 2020; Keyes, 2018).

To address these challenges, we introduce Gender Mistranslations Test Set (MiTens); a new dataset with 13 evaluation sets, including 26 languages (Table 1). We address challenges with contamination by creating targeted synthetic

datasets, releasing provenance of mined datasets, and marking dataset files with canaries (Srivastava et al., 2023). We address challenges with evaluation methods by precisely targeting specific error patterns, many of which can be scored automatically with simple heuristics. We additionally release evaluation sets for translating out of English, for use with human evaluation protocols similar to Anil et al. (2023). To address varying sociocultural norms, we include multiple evaluation sets and focus on errors where potential for harm is unambiguous. Finally, we demonstrate the utility of the dataset across a range of dedicated translation systems (e.g., NLLB, Team et al., 2022) and foundation models (e.g., GPT-4).

text[[115, 857, 486, 918], [511, 327, 881, 374]] We note that some languages we target such as Lingala have few existing evaluation resources. The evaluation sets we release can be expanded in future work (e.g., increasing diversity of source passages, more counterfactual variations). We also leave important challenges with mistranslation of non-binary gender expressions to future work.

Table 1: Languages included, grouped by level of digital resources, together with the number of examples in each group for translation into and out of English.

High	Mid	Low	Very low
Arabic	Finnish	Amharic	Assamese
Chinese	Indonesian	Bengali	Bhojpuri
French	Polish	Czech	Lingala
German	Telugu	Farsi	Luganda
Hindi	Turkish	Maithili	
Italian		Oromo	
Japanese			
Portuguese			
Russian			
Spanish			
#	2,252	488	784

2 Dataset

In order to precisely target different constructions and languages, and to enable fine-grained disaggregated evaluation, MiTTTeN5 contains multiple evaluation sets (Table 2). Evaluation sets target potential harm when translating into English ("2en"), or when translating from English into another language ("2xx"). To enable automated evaluation, all 2en evaluation sets are constructed so that the source language input contains only a single gendered entity. This enables automated scoring of English translation by scanning for the expression of grammatical gender in personal pronouns. Each data point contains around 1-10 sentences per source passage, and additionally includes a reference translation, with more details in the data card (Pushkarna et al., 2022). Evaluation sets are designed to pinpoint areas for improvement, rather than to exhaustively evaluate performance across all possible

source passages in each language.

2.1 Gender Sets

The Gender Sets evaluation set was built from error analysis in publicly available translation systems. The linguistic phenomena targeted include co-reference (Polish "Mój przyjaciel jest piosenkarszem, ale kompletnie bez talentu" to English "My friend is a singer but he is not talented at all"), gender agreement (Spanish "Mario trabaja jako empleado doméstico. Casi no pasa tiempo en su casa..." to English "Mario works as a housekeeper. He rarely spends time at home."), and gender-specific words (English "I went to my mother's house yesterday. She is British." to French "Je suis allé chez ma mère hier. Elle est britannique.").

Examples targeting co-reference were created using a mix of handwritten and synthetic methods. Examples targeting gender agreement were created from three sources: adapted from Translated Wikipedia Biographies (Stella, 2021), sourced from public news websites, or created synthetically. Examples targeting gender-specific words were created synthetically. Professional translators were used in creating reference translations. In total, this consists of 1,888 2xx data points. To enable automated evaluation for all 2en evaluation sets, we additionally filter those examples down to 630 2en data points. Filtering removes source passages with more than one English gender pronoun, and languages like Bengali that do not encode gender information in pronouns (this evaluation set only).

Table 2: Evaluation sets in MiTTTeN5.

Eval set	Subset	#
2xx: Translating out of English		
Gender Sets	coref:coreference	592
Gender Sets	coref:syntheticS	224
Gender Sets	gender_agreement:contextualS	496
Gender Sets	gender_agreement:news	192
Gender Sets	gender_agreement:wiki	256
Gender Sets	gender_specificS	128
2en: Translating into English		
Gender Sets	coref:coreference	180
Gender Sets	coref:syntheticS	210
Gender Sets	gender_agreement:contextualS	120
Gender Sets	gender_specificS	120
Late binding	late_binding	252
Enc in nouns	nouns_then_pronouns	222
SynthBio	synthbioS	640

2.2 SynthBio

The SynthBio evaluation set is mined from a subset of Yuan et al. (2022), which consists of synthetically generated English biography passages with multiple sentences. Using syn-

thetic data avoids potential data contamination from sources like Translated Wikipedia Biographies (Stella, 2021), which language models may have seen during pretraining. We filter SynthBio to only include passages encoding a single gendered entity with binary pronouns, then take a stratified sample based on English gender pronouns, and finally create pairs for a subset of languages using machine translation. This consists of 640 examples targeting translation into English. These passages often require gender information to be translated correctly across multiple sentences, and are longer passages. An example Thai to English reference translation is:

Suzanne Abamu was a Congolese feminist theologian, professor, and activist. Abamu was born on April 12, 1933 in Dekole, Republic of the Congo. She attended the University of Sorbonne Paris. She died on February 22, 2012 in Paris due to renal failure. She is buried in Cimetiere du Montparnasse in Paris. She is the daughter of Maria Abamu and Augustin Abamu. Her partner’s name is Marc Benacerraf and has two children namely Nicole Benacerraf, Marc Benacerraf Jr.

2.3 Late binding

The Late binding evaluation set was created from error analysis on translation errors in Gender Sets. It targets passages in Spanish where the gender information is only encoded later in the source passage, but where an English translation would require expression of gender early in the translation. For example in Spanish “Vino de inmediato cuando se enteró porque es una buena bibliotecaria” does not encode gender information until the end of the sentence, but in an English translation gender information would come early in “She came right away when she found out because she is a good librarian.” This evaluation set uses a mix of nouns for family names as well as a subset of nouns from Winogender (Rudinger et al., 2018), and consists of 252 examples targeting translation into English, including counterfactual passages.

2.4 Encoded in nouns

The Encoded in nouns evaluation set targets languages like Finnish that don’t encode gender information in personal pronouns but do encode gender information lexically through the choice of noun word (e.g., *isa* or *aiti*). This consists of 222 handcrafted examples targeting translation into English, with counterfactual passages that vary only by gender. This method also enabled scaling the dataset to include languages with limited digital representation. An example from the evaluation set in Oromo is “Saaraan akkoo kooti. Qoosaa ishee baay’een jaalladha.” with a reference translation of “Sarah is my aunt. I really like her jokes.”

3 Evaluation

MiTtenS can be used in evaluation for external audits of a deployed system, during model development, or monitoring during training. Here, we demonstrate using the dataset for

automated evaluation of 2en translation with a range of systems (details for reproducing are in Appendix A). For an 2xx human evaluation protocol see Anil et al. (2023). We leave demonstration of LLM-based evaluation (Zheng et al., 2023) for future work.

Evaluation results are shown in Figure 2, and we highlight specific areas of improvement for each system with disaggregated analysis by language and evaluation set in Table 3. Disaggregated analysis with precise evaluation data enables targeted improvements, and scales as additional evaluation sets are added over time. Even though systems show relatively high overall accuracy, in Figure 2 all systems perform worse on passages that require translation to “she” as compared to “he”, which may be related to patterns of representation in training datasets (Chowdhery et al., 2022). Performance in Table 3 is often worst on Encoded in nouns or Late binding evaluation sets. Surprisingly, we see areas of weakness even in high resource languages such as Spanish, and different areas of weakness in the same model families. There is no clear pattern to which languages are most challenging across systems, demonstrating the importance of empirical evaluations, and that MiTTenS can be used to pinpoint areas for targeted improvement.

IMAGE NOT PROVIDED

Figure 2: Evaluation results using automated evaluation when translating into English. Gemini and PaLM 2 systems perform best when considering worst-case performance, and GPT4 is within 5 percentage points.

Table 3: Automated evaluation results for translation into English.

Family	Model	Overall accuracy	Weakest language
NLLB	nllb-200-distilled-600M	98.0%	Bengali
GPT 4	gpt-4-1106-preview	99.1%	Lingala
GPT 3.5	gpt-3.5-turbo-1106	95.9%	Amharic
Gemini	gemini-pro	97.8%	Spanish
PaLM 2	text-bison-001	99.0%	Indonesian
PaLM 2	text-bison-32k	98.4%	Hindi
Mistral	Mistral-7B-Instruct-v0.1	92.7%	Lingala

4 Conclusion

We release MiTTenS, a dataset for measuring gender mistranslation harms with 13 evaluation sets that covers 26 languages. This dataset makes progress towards more precisely measuring potential harms and scaling evaluation to more languages. We address challenges with contamination and scoring methods amidst evolving sociocultural norms.

Future research should measure gender mistranslation in direct translation, expand automated evaluation methods, and to investigate how increasingly capable foundation models

might enable interactive or multiple alternative translations. More work is also needed to develop language technologies that produce accurate and faithful representations of non-binary people across all languages.

Limitations

For gender-related errors in translation systems, evaluations do not consider differential harms to people related to expressing non-binary gender identities (Keyes, 2018; Dev et al., 2021; Lauscher et al., 2023), or consider contested perspectives on pronouns across languages and cultures (Lee, 2019). Moreover, while gender agreement into English is amenable to automatic evaluation, evaluation of gender agreement out of English remains challenging and time-intensive. This dataset does not include examples for direct translation between languages beyond English, and it includes only a relatively small number of source passages. This dataset is not representative of the full range of human language and all passages that could be translated, which limits the comprehensiveness of evaluation results. This work is focused on translation when the gender information is unambiguously encoded in the source passage, and when there is a clear correct translation. Interpreting speaker or user intent in ambiguous contexts is a separate important class of evaluations with prior work, but one that this paper does not address. Finally, we note that this work focuses on only a subset of potential risks (Weidinger et al., 2021), and that our evaluations focus on model outputs without considering the wider sociotechnical context in which translation systems and foundation models exist (Weidinger et al., 2023; Shelby et al., 2023).

Ethical Considerations

This work aims to contribute to society and to human well-being by creating a new dataset and demonstrating how it can be used to measure some potential harms in translation systems. Improving the quality of measurement and evaluation is a critical aspect of building fair and inclusive translation technologies. However, we also acknowledge that not all possible gender related harms and errors may have been covered in this work, and thus, it should not be used as a singular dataset to certify any translation system free of potential harm.

In particular, this dataset is not able to cover non binary gendered pronouns and terms. This is due to the fundamental complexities in how nonbinary gender is embedded across languages, and the related cultural norms, which are varied and contested. Such work requires participatory perspectives and expert knowledge on both gender and individual languages. Gender mistranslations in these situations can result in misgendering harms that are especially salient and need to be studied deeply and with community engaged methods. Our dataset should not be used to measure this harm.

Earlier drafts of this paper used the term "misgendering" and we have revised our language in this draft thanks to

thoughtful reviewer feedback. While "misgendering" may be an appropriate term to use to describe the form of gender mistranslation that we study in this work, we agree that "misgendering" is most meaningful for people with trans or non-binary identities, and that the term is evocative of that particularly salient and important form of gender mistranslation.

We thank Marie Pellat, Orhan Firat, Kellie Webster, Kathy Meier-Hellstern, Erin van Liemt, Mark Diaz, and Amber Ebina for their input, feedback, and advice.

References

- [1] Syeda Nahida Akter, Zichun Yu, Aashiq Muhamed, Tianyue Ou, Alex Bauerle, Angel Alexander Cabrera, Krish Dholakia, Chenyan Xiong, and Graham Neubig. 2023. An in-depth look at gemini’s language abilities.
- [2] Bashar Alhafni, Nizar Habash, and Houda Bouamor. 2022. The Arabic parallel gender corpus 2.0: Extensions and analyses. In *Proceedings of the Thirteenth Language Resources and Evaluation Conference*, pages 1870-1884, Marseille, France. European Language Resources Association.
- [3] Rohan Anil, Andrew M. Dai, Orhan Firat, Melvin Johnson, Dmitry Lepikhin, Alexandre Passos, Siamak Shakeri, Emanuel Taropa, Paige Bailey, Zhifeng Chen, Eric Chu, Jonathan H. Clark, Laurent El Shafey, Yanping Huang, Kathy Meier-Hellstern, Gaurav Mishra, Erica Moreira, Mark Omernick, Kevin Robinson, Sebastian Ruder, Yi Tay, Kefan Xiao, Yuanzhong Xu, Yujing Zhang, Gustavo Hernandez Abrego, Junwhan Ahn, Jacob Austin, Paul Barham, Jan Botha, James Bradbury, Siddhartha Brahma, Kevin Brooks, Michele Catasta, Yong Cheng, Colin Cherry, Christopher A. Choquette-Choo, Aakanksha Chowdhery, Clement Crepy, Shachi Dave, Mostafa Dehghani, Sunipa Dev, Jacob Devlin, Mark Diaz, Nan Du, Ethan Dyer, Vlad Feinberg, Fangxiaoyu Feng, Vlad Fienber, Markus Freitag, Xavier Garcia, Sebastian Gehrmann, Lucas Gonzalez, Guy Gur-Ari, Steven Hand, Hadi Hashemi, Le Hou, Joshua Howland, Andrea Hu, Jeffrey Hui, Jeremy Hurwitz, Michael Isard, Abe Ittycheriah, Matthew Jagielski, Wenhao Jia, Kathleen Kenealy, Maxim Krikun, Sneha Kudugunta, Chang Lan, Katherine Lee, Benjamin Lee, Eric Li, Music Li, Wei Li, YaGuang Li, Jian Li, Hyeontaek Lim, and others. 2023. [MISSING].
- [4] Luisa Bentivogli, Beatrice Savoldi, and others. 2020. Gender in danger? evaluating speech translation technology on the MuST-SHE corpus. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 6923–6933, Online. Association for Computational Linguistics.
- [5] Su Lin Blodgett, Solon Barocas, Hal Daume III, and Hanna Wallach. 2020. Language (technology) is power:

A critical survey of "bias" in NLP. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 5454–5476, Online. Association for Computational Linguistics.

- [6] Su Lin Blodgett, Gilsinia Lopez, Alexandra Olteanu, Robert Sim, and Hanna Wallach. 2021. Stereotyping Norwegian salmon: An inventory of pitfalls in fairness benchmark datasets. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 1004–1015, Online. Association for Computational Linguistics.
- [7] Yang Trista Cao and Hal Daume III. 2020. Toward gender-inclusive coreference resolution. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 4568–4595, Online. Association for Computational Linguistics.
- [8] Won Ik Cho, Ji Won Kim, Seok Min Kim, and Nam Soo Kim. 2019. On measuring gender bias in translation of gender-neutral pronouns. In *Proceedings of the First Workshop on Gender Bias in Natural Language Processing*, pages 173–181, Florence, Italy. Association for Computational Linguistics.
- [9] Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, Parker Schuh, Kensen Shi, Sasha Tsvyashchenko, Joshua Maynez, Abhishek Rao, Parker Barnes, Yi Tay, Noam Shazeer, Vinodkumar Prabhakaran, Emily Reif, Nan Du, Ben Hutchinson, Reiner Pope, James Bradbury, Jacob Austin, Michael Isard, Guy Gur-Ari, Pengcheng Yin, Toju Duke, Anselm Levskaya, Sanjay Ghemawat, Sunipa Dev, Henryk Michalewski, Xavier Garcia, Vedant Misra, Kevin Robinson, Liam Fedus, Denny Zhou, Daphne Ippolito, David Luan, Hyeontaek Lim, Barret Zoph, Alexander Spiridonov, Ryan Sepassi, David Dohan, Shivani Agrawal, Mark Omernick, Andrew M. Dai, Thanumalayan Sankaranarayanan Pillai, Marie Pellat, Aitor Lewkowycz, Erica Moreira, Rewon Child, Oleksandr Polozov, Katherine Lee, Zongwei Zhou, Xuezhi Wang, Brennan Saeta, Mark Diaz, Orhan Firat, Michele Catasta, Jason Wei, Kathy Meier-Hellstern, Douglas Eck, Jeff Dean, Slav Petrov, and Noah Fiedel. 2022. Palm: Scaling language modeling with pathways.
- [10] Hyung Won Chung, Le Hou, Shayne Longpre, Barret Zoph, Yi Tay, William Fedus, Yunxuan Li, Xuezhi Wang, Mostafa Dehghani, Siddhartha Brahma, Albert Webson, Shixiang Shane Gu, Zhuyun Dai, Mirac Suzgun, Xinyun Chen, Aakanksha Chowdhery, Alex Castro-Ros, Marie Pellat, Kevin Robinson, Dasha Valter, Sharan Narang, Gaurav Mishra, Adams Yu, Vincent Zhao, Yanping Huang, Andrew Dai, Hongkun Yu, Slav Petrov, Ed H. Chi, Jeff Dean, Jacob Devlin, Adam Roberts, Denny Zhou, Quoc V. Le, and Jason Wei. 2022. Scaling instruction-finetuned language models.
- [11] Marta R Costa-jussa. 2019. An analysis of gender bias studies in natural language processing. *Nature Machine Intelligence*, 1(11):495–496.
- [12] Marta R Costa-jussa, Pierre Andrews, Eric Smith, Prangthip Hansanti, Christophe Ropers, Elahe Kalbassi, Cynthia Gao, Daniel Licht, and Carleigh Wood. 2023. Multilingual holistic bias: Extending descriptors and patterns to unveil demographic biases in languages at scale. *arXiv preprint arXiv:2305.13198*.
- [13] Marta R. Costa-jussa, Eric Smith, Christophe Ropers, Daniel Licht, Jean Maillard, Javier Ferrando, and Carlos Escolano. 2023. Toxicity in multilingual machine translation at scale.
- [14] Sunipa Dev, Masoud Monajatipoor, Anaelia Ovalle, Arjun Subramonian, Jeff M Phillips, and Kai-Wei Chang. 2021. Harms of gender exclusivity and challenges in non-binary representation in language technologies.
- [15] Gemini Team Google. 2023. Gemini: A family of highly capable multimodal models. *arXiv preprint arXiv:2312.11805*.
- [16] Seraphina Goldfarb-Tarrant, Rebecca Marchant, Ricardo Munoz Sanchez, Mugdha Pandya, and Adam Lopez. 2021. Intrinsic bias metrics do not correlate with application bias.
- [17] Tamanna Hossain, Sunipa Dev, and Sameer Singh. 2023. Misgendered: Limits of large language models in understanding pronouns.
- [18] Melvin Johnson. 2018. Providing gender-specific translations in google translate.
- [19] Melvin Johnson. 2020. A scalable approach to reducing gender bias in google translate.
- [20] Yennie Jun. 2023. Lost in dall-e 3 translation.
- [21] Os Keyes. 2018. The misgendering machines: Trans/hci implications of automatic gender recognition. *Proc. ACM Hum.-Comput. Interact.*, 2(CSCW).
- [22] Douwe Kiela, Max Bartolo, Yixin Nie, Divyansh Kaushik, Atticus Geiger, Zhengxuan Wu, Bertie Vidgen, Grusha Prasad, Amanpreet Singh, Pratik Ringshia, Zhiyi Ma, Tristan Thrush, Sebastian Riedel, Zeerak Waseem, Pontus Stenetorp, Robin Jia, Mohit Bansal, Christopher Potts, and Adina Williams. 2021. Dynabench: Rethinking benchmarking in nlp.
- [23] Jack Krawczyk. 2023. Bard’s latest update: more features, languages and countries.

- [24] Anne Lauscher, Debora Nozza, Ehm Miltersen, Archie Crowley, and Dirk Hovy. 2023. What about “em”? how commercial machine translation fails to handle (neo-)pronouns. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 377-392, Toronto, Canada. Association for Computational Linguistics.
- [25] Chelsea Lee. 2019. Welcome, singular “they”. <https://apastyle.apa.org/blog/singular-they>. Accessed: 2022-11-18.
- [26] OpenAI, Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Alstenschmidt, Sam Altman, Shyamal Anadkat, Red Avila, Igor Babuschkin, Suchir Balaji, Valerie Balcom, Paul Baltescu, Haiming Bao, Mo Bavarian, Jeff Belgum, Irwan Bello, Jake Berdine, Gabriel Bernadett-Shapiro, Christopher Berner, Lenny Bogdonoff, Oleg Boiko, Madelaine Boyd, Aina-Luisa Brakman, Greg Brockman, Tim Brooks, Miles Brundage, Kevin Button, Trevor Cai, Rosie Campbell, Andrew Cann, Brittany Carey, Chelsea Carlson, Rory Carmichael, Brooke Chan, Che Chang, Fotis Chantzis, Derek Chen, Sully Chen, Ruby Chen, Jason Chen, Mark Chen, Ben Chess, Chester Cho, Casey Chu, Hyung Won Chung, Dave Cummings, Jeremiah Currier, Yunxing Dai, Cory De-careaux, Thomas Degry, Noah Deutsch, Damien Deville, Arka Dhar, David Dohan, Steve Dowling, Sheila Dunning, Adrien Ecoffet, Atty Eleti, Tyna Eloundou, David Farhi, Liam Fedus, Niko Felix, Simon Posada Fishman, Juston Forte, Isabella Fulford, Leo Gao, Elie Georges, Christian Gibson, Vik Goel, Tarun Gogineni, Gabriel Goh, Rapha Gontijo-Lopes, Jonathan Gordon, Morgan Grafstein, Scott Gray, Ryan Greene, Joshua Gross, Shixiang Shane Gu, Yufei Guo, Chris Hallacy, Jesse Han, Jeff Harris, Yuchen He, Mike Heaton, Johannes Heidecke, Chris Hesse, Alan Hickey, Wade Hickey, Peter Hoeschele, Brandon Houghton, Kenny Hsu, Shengli Hu, Xin Hu, Joost Huizinga, Shantanu Jain, Shawn Jain, Joanne Jang, Angela Jiang, Roger Jiang, Haozhun Jin, Denny Jin, Shino Jomoto, Billie Jonn, Heewoo Jun, Tomer Kaftan, Łukasz Kaiser, Ali Kamali, Ingmar Kanitscheider, Nitish Shirish Keskar, Tabarak Khan, Logan Kilpatrick, Jong Wook Kim, Christina Kim, Yongjik Kim, Hendrik Kirchner, Jamie Kiros, Matt Knight, Daniel Kokotajlo, Łukasz Kondraciuk, Andrew Kondrich, Aris Konstantinidis, Kyle Kosic, Gretchen Krueger, Vishal Kuo, Michael Lampe, Ikai Lan, Teddy Lee, Jan Leike, Jade Leung, Daniel Levy, Chak Ming Li, Rachel Lim, Molly Lin, Stephanie Lin, Mateusz Litwin, Theresa Lopez, Ryan Lowe, Patricia Lue, Anna Makanju, Kim Malfacini, Sam Manning, Todor Markov, Yaniv Markovski, Bianca Martin, Katie Mayer, Andrew Mayne, Bob McGrew, Scott Mayer McKinney, Christine McLeavey, Paul McMillan, Jake McNeil, David Medina, Aalok Mehta, Jacob Menick, Luke Metz, Andrey Mishchenko, Pamela Mishkin, Vinnie Monaco, Evan Morikawa, Daniel Mossing, Tong Mu, Mira Murati, Oleg Murk, David Mély, Ashvin Nair, Reiichiro Nakano, Rajeev Nayak, Arvind Neelakantan, Richard Ngo, Hyeonwoo Noh, Long Ouyang, Cullen O’Keefe, Jakub Pachocki, Alex Paino, Joe Palermo, Ashley Pantuliano, Giambattista Parascandolo, Joel Parish, Emy Parparita, Alex Passos, Mikhail Pavlov, Andrew Peng, Adam Perelman, Filipe de Avila Belbute Peres, Michael Petrov, Henrique Ponde de Oliveira Pinto, Michael, Pokorný, Michelle Pokrass, Vitchyr Pong, Tolly Powell, Alethea Power, Boris Power, Elizabeth Proehl, Raul Puri, Alec Radford, Jack Rae, Aditya Ramesh, Cameron Raymond, Francis Real, Kendra Rimbach, Carl Ross, Bob Rotsted, Henri Roussez, Nick Ryder, Mario Saltarelli, Ted Sanders, Shibani Santurkar, Girish Sastry, Heather Schmidt, David Schnurr, John Schulman, Daniel Selsam, Kyla Sheppard, Toki Sherbakov, Jessica Shieh, Sarah Shoker, Pranav Shyam, Szymon Sidor, Eric Sigler, Maddie Simens, Jordan Sitkin, Katarina Slama, Ian Sohl, Benjamin Sokolowsky, Yang Song, Natalie Staudacher, Felipe Petroski Such, Natalie Summers, Ilya Sutskever, Jie Tang, Nikolas Tezak, Madeleine Thompson, Phil Tillet, Amin Tootoonchian, Elizabeth Tseng, Preston Tuggle, Nick Turley, Jerry Tworek, Juan Felipe Cerón Uribe, Andrea Vallone, Arun Vijayvergiya, Chelsea Voss, Carroll Wainwright, Justin Jay Wang, Alvin Wang, Ben Wang, Jonathan Ward, Jason Wei, CJ Weinmann, Akila Welihinda, Peter Welinder, Jiayi Weng, Lilian Weng, Matt Wiethoff, Dave Willner, Clemens Winter, Samuel Wolrich, Hannah Wong, Lauren Workman, Sherwin Wu, Jeff Wu, Michael Wu, Kai Xiao, Tao Xu, Sarah Yoo, Kevin Yu, Qiming Yuan, Wojciech Zaremba, Rowan Zellers, Chong Zhang, Marvin Zhang, Shengjia Zhao, Tianhao Zheng, Juntang Zhuang, William Zhuk, and Barret Zoph. 2023. Gpt4 technical report.
- [27] Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Christiano, Jan Leike, and Ryan Lowe. 2022. Training language models to follow instructions with human feedback.
- [28] Mahima Pushkarna, Andrew Zaldivar, and Oddur Kjar-tansson. 2022. Data cards: Purposeful and transparent dataset documentation for responsible ai.
- [29] Adithya Renduchintala, Denise Diaz, Kenneth Heafield, Xian Li, and Mona Diab. 2021. Gender bias amplification during speed-quality optimization in neural machine translation. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Nat-*

ural Language Processing (Volume 2: Short Papers), pages 99-109, Online. Association for Computational Linguistics.

- [30] Rachel Rudinger, Jason Naradowsky, Brian Leonard, and Benjamin Van Durme. 2018. Gender bias in coreference resolution. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 2 (Short Papers)*, pages 8-14, New Orleans, Louisiana. Association for Computational Linguistics.
- [31] Beatrice Savoldi, Marco Gaido, Luisa Bentivogli, Matteo Negri, and Marco Turchi. 2021. Gender bias in machine translation. *Transactions of the Association for Computational Linguistics*, 9:845-874.
- [32] Renee Shelby, Shalaleh Rismani, Kathryn Henne, AJung Moon, Negar Rostamzadeh, Paul Nicholas, N'Mah Yilla, Jess Gallegos, Andrew Smart, Emilio Garcia, and Gurleen Virk. 2023. Sociotechnical harms of algorithmic systems: Scoping a taxonomy for harm reduction.
- [33] Pushdeep Singh. 2023a. Don't overlook the grammatical gender: Bias evaluation for hindi-english machine translation.
- [34] Pushdeep Singh. 2023b. Gender inflected or bias inflected: On using grammatical gender cues for bias evaluation in machine translation.
- [35] AaroHi Srivastava, Abhinav Rastogi, Abhishek Rao, Abu Awal Md Shoeb, Abubakar Abid, Adam Fisch, Adam R. Brown, Adam Santoro, Aditya Gupta, Adria Garriga-Alonso, Agnieszka Kluska, Aitor Lewkowycz, Akshat Agarwal, Alethea Power, Alex Ray, Alex Warstadt, Alexander W. Kocurek, Ali Safaya, Ali Tazarv, Alice Xiang, Alicia Parrish, Allen Nie, Aman Hussain, Amanda AskeIl, Amanda Dsouza, Ambrose Slone, Ameet Rahane, Anantharaman S. Iyer, Anders Andreassen, Andrea Madotto, Andrea Santilli, Andreas Stuhlmüller, Andrew Dai, Andrew La, Andrew Lampinen, Andy Zou, Angela Jiang, Angelica Chen, Anh Vuong, Animesh Gupta, Anna Gottardi, Antonio Norelli, Anu Venkatesh, Arash Gholami-davoodi, Arfa Tabassum, Arul Menezes, Arun Kirubakaran, Asher Mullokandov, Ashish Sabharwal, Austin Herrick, Avia Efrat, Aykut Erdem, Ayla Karakas, B. Ryan Roberts, Bao Sheng Loe, Barret Zoph, Bartłomiej Bojanowski, Batuhan Özyurt, Behnam Hedayatnia, Behnam Neyshabur, Benjamin Inden, Benno Stein, Berk Ekmekci, Bill Yuchen Lin, Blake Howald, Bryan Orinon, Cameron Diao, Cameron Dour, Catherine Stinson, Cedrick Argueta, César Ferri Ramírez, Chandan Singh, Charles Rathkopf, Chenlin Meng, Chitta Baral, Chiyu Wu, Chris Callison-Burch, Chris Waites, Christian Voigt, Christopher D. Manning, Christopher Potts, Cindy Ramirez, Clara E. Rivera, Clemencia Siro, Colin Raffel, Courtney Ashcraft, Cristina Garbacea, Damien Sileo, Dan Garrette, Dan Hendrycks, Dan Kilman, Dan Roth, Daniel Freeman, Daniel Khashabi, Daniel Levy, Daniel Moseguí González, Danielle Perszyk, Danny Hernandez, Danqi Chen, Daphne Ippolito, Dar Gilboa, David Dohan, David Drakard, David Jurgens, Debajyoti Datta, Deep Ganguli, Denis Emelin, Denis Kleyko, Deniz Yuret, Derek Chen, Derek Tam, Dieuwke Hupkes, Diganta Misra, Dilyar Buzan, Dimitri Coelho Mollo, Diyi Yang, Dong-Ho Lee, Dylan Schrader, Ekaterina Shutova, Ekin Dogus Cubuk, Elad Segal, Eleanor Hagerman, Elizabeth Barnes, Elizabeth Donoway, Ellie Pavlick, Emanuele Rodola, Emma Lam, Eric Chu, Eric Tang, Erkut Erdem, Ernie Chang, Ethan A. Chi, Ethan Dyer, Ethan Jerzak, Ethan Kim, Eunice Engefu Manyasi, Evgenii Zheltonozhskii, Fanyue Xia, Fatemeh Siar, Fernando Martínez-Plumed, Francesca Happé, Francois Chollet, Frieda Rong, Gaurav Mishra, Genta Indra Winata, Gerard de Melo, Germán Kruszewski, Giambattista Parascandolo, Giorgio Mariani, Gloria Wang, Gonzalo Jaimovitch-López, Gregor Betz, Guy Gur-Ari, Hana Galijasevic, Hannah Kim, Hannah Rashkin, Hannaneh Hajishirzi, Harsh Mehta, Hayden Bogar, Henry Shevlin, Hinrich Schütze, Hiromu Yakura, Hongming Zhang, Hugh Mee Wong, Ian Ng, Isaac Noble, Jaap Jumelet, Jack Geissinger, Jackson Kernion, Jacob Hilton, Jaehoon Lee, Jaime Fernández Fisac, James B. Simon, James Koppel, James Zheng, James Zou, Jan Kocof, Jana Thompson, Janelle Wingfield, Jared Kaplan, Jaroma Radom, Jascha Sohl-Dickstein, Jason Phang, Jason Wei, Jason Yosinski, Jekaterina Novikova, Jelle Bosscher, Jennifer Marsh, Jeremy Kim, Jeroen Taal, Jesse Engel, Jesujoba Alabi, Jiacheng Xu, Jiaming Song, Jillian Tang, Joan Waweru, John Burden, John Miller, John U. Balis, Jonathan Batchelder, Jonathan Berant, Jörg Frohberg, Jos Rozen, Jose Hernandez-Orallo, Joseph Boudeman, Joseph Guerr, Joseph Jones, Joshua B. Tenenbaum, Joshua S. Rule, Joyce Chua, Kamil Kanclerz, Karen Livescu, Karl Krauth, Karthik Gopalakrishnan, Katerina Ignatyeva, Katja Markert, Kaustubh D. Dhole, Kevin Gimpel, Kevin Omondi, Kory Mathewson, Kristen Chiafullo, Ksenia Shkaruta, Kumar Shridhar, Kyle McDonell, Kyle Richardson, Laria Reynolds, Leo Gao, Li Zhang, Liam Dugan, Lianhui Qin, Lidia Contreras-Ochando, Louis-Philippe Morency, Luca Moschella, Lucas Lam, Lucy Noble, Ludwig Schmidt, Luheng He, Luis Oliveros Colón, Luke Metz, Lütfi Kerem Şenel, Maarten Bosma, and others. 2023. [MISSING].
- [36] Karolina Stanczak and Isabelle Augenstein. 2021. A survey on gender bias in natural language processing.
- [37] Gabriel Stanovsky, Noah A. Smith, and Luke Zettlemoyer. 2019. Evaluating gender bias in machine trans-

lation. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 1679-1684, Florence, Italy. Association for Computational Linguistics.

All evaluation results are from December 2023. At the time of writing in June 2024, we note that the specific 'gemini-pro' system evaluated is no longer available.

- [38] Romina Stella. 2021. Translated wikipedia biographies.
- [39] Tao Sun, et al. 2019. [MISSING].
- [40] Eva Vanmassenhove, Christian Hardmeier, and Andy Way. 2018. Getting gender right in neural machine translation. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 3003-3008, Brussels, Belgium. Association for Computational Linguistics.
- [41] Laura Weidinger, John Mellor, Maribeth Rauh, Conor Griffin, Jonathan Uesato, Po-Sen Huang, Myra Cheng, Mia Glaese, Borja Balle, Atoosa Kasirzadeh, et al. 2021. Ethical and social risks of harm from language models.
- [42] Laura Weidinger, Maribeth Rauh, Nahema Marchal, Arianna Manzini, Lisa Anne Hendricks, Juan Mateos-Garcia, Stevie Bergman, Jackie Kay, Conor Griffin, Ben Bariach, Iason Gabriel, Verena Rieser, and William Isaac. 2023. Sociotechnical safety evaluation of generative ai systems.
- [43] Shuo Yang, Wei-Lin Chiang, Lianmin Zheng, Joseph E. Gonzalez, and Ion Stoica. 2023. Rethinking benchmark and contamination for language models with rephrased samples.
- [44] Zheng-Xin Yong, Cristina Menghini, and Stephen H. Bach. 2023. Low-resource languages jailbreak gpt4.
- [45] Ann Yuan, Daphne Ippolito, Vitaly Nikolaev, Chris Callison-Burch, Andy Coenen, and Sebastian Gehrmann. 2022. Synthbio: A case study in human-ai collaborative curation of text datasets.
- [46] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. Judging llm-as-a-judge with mt-bench and chatbot arena.

A Evaluation protocol details

GPT systems were queried with the OpenAI Python client, and PaLM 2 and Gemini systems with the Cloud Vertex Python SDK. Mistral was evaluated through a HuggingFace Endpoint. NLLB was run in local inference.

Foundation models were prompted with an instruction with greedy sampling (top-k=1 or temperature=0), using the instruction below, shown with an example prompt to translate a Turkish source passage into English.

Translate the following text from Turkish to English.

Turkish: Sarah bir aktris. Yakınlarda yaşıyor. English: