

SparseGrad: A Selective Method for Efficient Fine-tuning of MLP Layers

Viktoriia Chekalina^{1,2} Anna Rudenko^{1,2} Gleb Mezentsev^{1,2} Alexander Mikhalev² Alexander Pancher

¹Artificial Intelligence Research Institute

²Skolkovo Institute of Science and Technology

Abstract

The performance of Transformer models has been enhanced by increasing the number of parameters and the length of the processed text. Consequently, fine-tuning the entire model becomes a memory-intensive process. High-performance methods for parameter-efficient fine-tuning (PEFT) typically work with Attention blocks and often overlook MLP blocks, which contain about half of the model parameters. We propose a new selective PEFT method, namely SparseGrad, that performs well on MLP blocks. We transfer layer gradients to a space where only about 1% of the layer’s elements remain significant. By converting gradients into a sparse structure, we reduce the number of updated parameters. We apply SparseGrad to fine-tune BERT and RoBERTa for the NLU task and LLaMa-2 for the Question-Answering task. In these experiments, with identical memory requirements, our method outperforms LoRA and MeProp, robust popular state-of-the-art PEFT approaches.

1 Introduction

Due to the tendency to increase the size of transformer models with each new generation, we need efficient ways to fine-tune such models on downstream task data. The usual practice is fine-tuning a large pre-trained foundational model on a downstream task. The major problem that prevents efficient fine-tuning is a steady increase in the memory footprint. One of the best strategies is high-performance methods for parameter-efficient fine-tuning (PEFT). Typically, such methods as LoRA focus on attention blocks and do not consider dense MLP blocks. Since MLP blocks can take a significant fraction of the model parameters (see Table 1), we propose to focus instead on MLP blocks. We introduce a novel selective PEFT approach called SparseGrad. Our method is based on finding a special sparsification transformation that allows us to fine-tune about 1% of the dense MLP layer parameters and still show good performance in downstream tasks.

Table 1: Number of parameters for different layers in models based on the Transformer.

Blocks/Model	BERT _{base}		RoBERTa _{base}		LLaMa-2	
	Params	%	Params	%	Params	%
Full model	109 M	100%	125 M	100%	6.7 B	100%
MLP	57 M	52%	57 M	45%	4.3 B	64%
Embeddings	24 M	22%	40 M	32%	0.1 B	1%
Attention	28 M	25%	28 M	22%	2.1 B	31%

We validate our approach on BERT and RoBERTa models on GLUE benchmark and in both cases obtain results better than LoRA and MeProp methods. We also fine-tune LLaMa-2 7B on the OpenAssistant dataset and also achieve performance higher than LoRA and MeProp.

2 Related Work

In the last few years, many approaches to PEFT have appeared. ? distinguishes three types of methods: additive, reparametrization-based, and selective. In additive PEFT, small neural networks called adapters are added to the main model to steer the outputs of its modules ?. Adapters are trainable, therefore, the main model remains unchanged. ? adapt this approach to NLP. In reparametrization-based approaches low-rank representations of trainable parameters are used. For example, LoRA? parameterizes the weight update by a trainable low-rank matrix decomposition. In the original paper, LoRA is applied to self-attention modules, but not to MLP ones. In the selective methods, parts of the model or sets of the parameters are chosen for fine-tuning using some heuristics. Such methods include, for example, BitFit ? or MeProp?, where only top-k parameters are updated during backpropagation. The approach proposed in this paper is related to selective methods.

3 Method

Our aim is to reduce the amount of trainable parameters at the fine-tuning stage. Taking into account that fine-tuning data is restricted to a limited scope, we assume there is a basis where the weight gradient matrix is very close to being sparse. To identify this basis, we applied a decomposition technique to the stacked weight gradient matrices. As a result, we introduce a new PyTorch layer class, `SparseGradLinear`, which transitions weights to this sparse gradient space, accumulates gradients in sparse form, and enables the reverse transition back to the original space.

3.1 Preliminary Phase: Finding Transition Matrices

To obtain transition matrices, an initial procedure is necessary. During this, we perform n_{steps} of standard backpropagation by freezing the entire model and unfreezing only the linear layers in MLP blocks. We do it to obtain the set of weights gradient matrices $\frac{\partial L}{\partial W} \in \mathbb{R}^{D_{\text{in}} \times D_{\text{out}}}$. Stacking these matrices over n_{blocks} – the number of all blocks in the model – and over n_{steps} , we obtain a 3D tensor of size $D_{\text{in}} \times D_{\text{out}} \times (n_{\text{steps}} * n_{\text{blocks}})$.

Applying Higher Order SVD (HOSVD) ? to this tensor yields matrices $U \in \mathbb{R}^{D_{\text{in}} \times D_{\text{in}}}$, corresponding to the dimension D_{in} and $V^T \in \mathbb{R}^{D_{\text{out}} \times D_{\text{out}}}$, corresponding to D_{out} . In this way, we get two orthogonal transition matrices U, V^T which are shared across all blocks of the model. Multiplying the layer’s weight matrix on the left by U and on the right by V^T transforms it into a new space. In this transformed space, the gradient matrix exhibits greater sparsity compared to the original space. Examples of $\frac{\partial L}{\partial W^T}$ with and without transition to the new space are shown in Fig. 2.

3.2 Signal Propagation in SparseGradLinear Layer

Given a Transformer Linear layer with a weight matrix W^T , input activation X , and output $Y = XW^T$, we define the gradients of the output, input, and weights as $\frac{\partial L}{\partial Y}$, $\frac{\partial L}{\partial X}$, and $\frac{\partial L}{\partial W^T}$, respectively. To create the corresponding `SparseGradLinear` layer, we represent the weights in the U, V^T basis, such that the new weights are $\tilde{W}^T = UW^TV^T$. Since the modules following `SparseGradLinear` remain unchanged in both forward and backward passes, it is crucial to maintain consistency between outputs of the Original Linear Layer Y and the `SparseGradLinear` layer $\tilde{Y}$, as well as their input gradients $\frac{\partial L}{\partial X}$ and $\frac{\partial L}{\partial \tilde{X}}$.

Table 2 outlines these adjustments and illustrates the correspondence of variables in Torch Autograd for Linear and `SparseGradLinear` layers.

Thus, `SparseGradLinear` is equivalent to 3 linear layers: first with frozen weights U^T , defined by the HOSVD, second with trainable new weights $\tilde{W}^T = UW^TV^T$, third with frozen weights V , defined by the HOSVD. Fig. 1 shows the propagation of the signal in this structure.

Table 2: Correspondence of variables in Torch Autograd for a regular Linear layer and SparseGradLinear.

Variable / Layer	Linear	SparseGrad
Weights	W^T	$\tilde{W}^T = UW^TV^T$
Input	X	$\tilde{X} = XU^T$
Output	$Y = XW^T$	$\tilde{Y} = XU^T\tilde{W}^TV$
Grad Output	$\frac{\partial L}{\partial Y}$	$\frac{\partial L}{\partial \tilde{Y}}V^T$
Grad Input	$\frac{\partial L}{\partial Y}W$	$\frac{\partial L}{\partial \tilde{Y}}U^T\tilde{W}^TV$
Grad Weights	$\frac{\partial L}{\partial Y}^T X$	$\frac{\partial L}{\partial \tilde{Y}}^T XU^T$

3.3 Sparse-by-Dense Matrix Multiplication

We provide the SparseGradLinear class with updated Forward and Backward procedures. However, the addition of multiplications by U , V into them increased the execution time and affected peak memory in the training loop.

The sparsity of the gradient tensor $\frac{\partial L}{\partial W} = \frac{\partial L}{\partial \tilde{Y}}^T \tilde{X}$ results in some of the multipliers being sparse. We explore the structure of each component in this formula and figure out that $\frac{\partial L}{\partial \tilde{Y}}$ has a sparsity approximately equal to $\frac{\partial L}{\partial W}$. Histograms of the percent of its non-zero elements are presented in Fig. 3. It also shows that the sparsity is “strided” – most of the rows are completely filled with zeros. These rows can be excluded from the multiplication procedure, thus optimizing it.

More precisely, to multiply the sparse matrix $A \in \mathbb{R}^{b \times c}$ by a dense matrix $B \in \mathbb{R}^{c \times d}$ we select `rows` and `cols` – indices of rows and columns of A which contain nonzero elements and multiply as follows:

$$C = A(\text{rows}, :)(:, \text{cols})B(\text{cols}, :). \quad (1)$$

We employ C either for further multiplications, or convert it into COO format and send it to SparseAdam optimizer. Indexes in COO format are defined by restoring indexes of A :

$$C_{\text{coo}}(\text{rows}(k), \text{cols}(l)) = C(k, l). \quad (2)$$

As it is shown in Table 3, such procedure significantly speeds up the harnessing of SparseGradLinear.

4 Time and Memory Consumption per Training Iteration

We measure the peak memory allocated during training using the CUDA memory allocator statistics. Table 3 demonstrates this statistic on average for all GLUE datasets for the RoBERTa_{base} model. The comprehensive Tables 7 and 8, which outline metrics for each dataset separately, can be found in Appendix A. Among all methods, LoRA presents the most efficient memory usage, preserving 30% of the peak memory. SparseGrad, while using slightly more memory, still achieves a 20% savings. The increase in peak memory with SparseGrad is attributed to the maintenance of matrices U and V and their multiplication by the dense objects, such as Input X .

In terms of training time, LoRA demonstrates the fastest training, followed by SparseGrad, and then standard fine-tuning. Table 3 shows that Sparse-by-Dense multiplication saves approximately 12% memory, leading to an almost five-fold increase in speed.

5 Experiments

We conducted experiments on three transformer-based encoder models, BERT and RoBERTa base and large, on the GLUE ? benchmark, and the LLaMa-2 decoder model on the OpenAssistant Conversations

Table 3: Training speed and memory requirements averaged on the GLUE benchmark. The last two rows of the Table report the results for the SparseGradmethod with Sparse-by-Dense (SD) and Regular (Reg) matrix multiplication, respectively.

Method	Steps/Sec.	Memory, MB
Regular FT	4.11	1345
LoRA	4.7	944
SparseGradSD	4.3	1016
SparseGradReg	0.9	1210

corpus ?. We compared the fine-tuning of the full model (Regular FT scheme) with three PEFT methods, namely LoRA, MePropand SparseGrad, applied to MLP blocks. To harness LoRA, we use an official repository code. For the MePropmethod, we kept the largest elements in the $\frac{\partial L}{\partial W}$ matrix. The proposed SparseGradinvolves replacing layers in MLP blocks with its SparseGradLinear equivalents.

5.1 Natural Language Understanding with BERT and RoBERTa

We explore the acceptable sparsity level of the gradient matrices in the “sparse” space, $\frac{\partial L}{\partial W}$. By varying the number of remaining parameters in the Linear Layer from $100 \cdot 10^3$ to $18 \cdot 10^3$, we fine-tuned the model on the GLUE benchmark and identified the point at which performance begins to degrade. This occurs when the number of trainable parameters reaches 22×10^3 , corresponding to 1% of the total weights. Full experimental results can be found in Appendix C.

Guided by this heuristic, in our experiments we leave the top 1% of the largest elements and set the rest to zero. To deal with SparseGradients, we use the SparseAdam optimizer – the masked version of the Adam algorithm. The remaining model parameters are trained with the standard AdamW optimizer.

We fine-tune BERT, RoBERTa_{base} and RoBERTa_{large} ? using Regular FT, LoRA, MePropand SparseGradschemes for 20 epochs with early stopping for each task in the GLUE. We varied the batch size and learning rate using the Optuna framework ?. The learning rate ranged from 10^{-6} to 10^{-1} , and the batch size is selected from the set $\{8, 16, 32\}$. Optimal training parameters for each task are available in the Appendix D. In LoRAwe take the rank 10 for RoBERTa_{large} and rank 7 for BERT and RoBERTa_{base}. For SparseGradand MePropwe keep the same number of parameters – approximately 1% of each Linear layer.

The average scores for all GLUE tasks for BERT and RoBERTa_{base} are in Table 4; per-task results are placed in the Appendix B. Table 5 depicts the scores for the RoBERTa_{large} model.

Table 4: Average scores over the GLUE benchmark for BERT and RoBERTa_{base} models.

Model AVG	Method	#Trainable params	Model	MLP block
BERT	Regular FT	109 mln	3 mln	82.5
	LoRA	54 mln	0.03 mln	81.6
	SparseGrad	54 mln	0.03 mln	82.6
	MeProp	54 mln	0.03 mln	82.1
RoBERTa _{base}	Regular FT	125 mln	3 mln	84.2
	LoRA	68 mln	0.03 mln	83.1
	SparseGrad	68 mln	0.03 mln	83.6
	MeProp	68 mln	0.03 mln	82.5

Our results indicate that SparseGradoutperforms LoRAwith an equivalent number of trainable parameters across all models. For BERT, SparseGradeven exceeds the performance of Regular FT. This

Table 5: Comparative results of RoBERTa_{large} for 20-epoch task-specific fine-tuning.

Method SST2	#Trainable params	Model	MLP block	AVG	STSB	CoLA	MNLI	MRPC	
Regular FT	355 mln	4 mln	85.6	91.9±.4	67.1±2.3	90.8±.2	89.9±.3	92.9±.9	9
LoRA	168 mln	0.05 mln	83.7	92.1±.3	64.4±.8	90.7±.2	89.9±.3	93.2±.3	9
SparseGrad	168 mln	0.05 mln	85.4	92.4±.2	63.2±3.4	90.7±.2	90.5±.5	93.3±.5	9
MeProp	168 mln	0.05 mln	84.3	92.3±.1	63.7±1.1	90.4±.2	89.4±.9	92.5±.5	9

may be attributed to the changing basis of the weights in SparseGradacting as a form of regularization. Concerning MeProp, it provides weaker results than SparseGrad in all cases except the RoBERTa_{large} on CoLA. This could be explained by the fact that our approach first transforms the elements into a special “sparse” space, while MeProp operates on gradients in the original space. In the original space, the histogram of elements is flatter (see Fig. 2), which suggests that, with the same cut-off threshold, MeProp may remove more significant elements compared to SparseGrad.

5.2 Conversations with LLaMa-2

We apply the SparseGrad method to fine-tune LLaMa-2 7B ? model on the OpenAssistant conversational dataset ?. Fine-tuning was performed on a single GPU NVIDIA A40 during 1 epoch with learning rate 9×10^{-4} . For Regular FT, we unfroze up_proj and down_proj layers in the MLP modules with a block index divisible by 3 (0, 3, 6, ...). We apply LoRA with rank 32 to the selected blocks, leaving the rest of the model untrainable. In the SparseGrad and MeProp methods, we also consider selected MLP modules in the transformer and leave $\approx 100,000$ (0.2%) nonzero elements in the gradient matrix. For LLaMA-2, we conducted a similar ablation study as we did for BERT and RoBERTa. We varied the number of remaining parameters in the MLP block and identified the point where the model’s performance began to decline.

We validate obtained models on the question set MT-Bench Inf from Inflection-Benchmarks ?. We followed the guidelines outlined in this work, called “Single Protocol” or “Single Answer Grading”. We got the answers by using the FastChat platform¹ and then evaluating them using GPT-4. GPT-4 rates the answers on a scale of 1 to 10, with the evaluation prompt taken from ?.

The resulting losses and average GPT-4 scores are presented in Table 6. While the models perform similarly overall, SparseGrad slightly outperforms LoRA, MeProp, and regular fine-tuning. Examples of responses to Inflection-Benchmark samples are provided in Appendix E. These examples illustrate that, although all models produce good answers, the LoRA-trained model occasionally overlooks important nuances. In the examples given, it fails to recognize that presentations can be stressful for introverts or that hierarchy plays a significant role in Japanese corporate culture.

Table 6: Comparative results for LLaMa-2 on the OpenAssistant dataset.

Method	#Train params	Valid Loss	I-Bench Score
Regular FT	22%	1.250 ±0.03	4.407
LoRA	0.5%	1.249 ±0.05	5.025
SparseGrad	0.5%	1.247 ±0.03	5.132
MeProp	0.5%	1.259 ±0.04	4.261

¹<https://github.com/lm-sys/FastChat>

6 Conclusion

We propose a new selective PEFT method called SparseGrad, which identifies a space where the gradients exhibit a sparse structure and updates only its significant part. SparseGrad is validated through experiments conducted on the BERT, RoBERTa and LLaMa-2 models, demonstrating its superiority over the additive LoRA and selective MeProp methods.

Leveraging the sparsity property significantly accelerated the calculations in SparseGrad. Our method runs faster than standard fine-tuning but slower than LoRA, while yielding better performance than LoRA; the same trend applies to memory usage. In summary, our method serves as an alternative to LoRA in situations where the performance of the final model takes precedence over the execution time. The source code as well as links to pretrained models are available at repository.²

7 Acknowledgements

The work was supported by the Analytical center under the RF Government (subsidy agreement 000000D730321P5Q000 Grant No. 70-2021-00145 02.11.2021).

8 Limitations

The main limitation of our method is the additional memory requirements during the Preliminary Phase. The extra memory is assessed as follows: we need to unfreeze the MLP layers, which hold approximately half of the training parameters in Transformers (see Table 1), store and decompose a large tensor. For instance, 30 steps in the preliminary phase result in a tensor of approximately 276 MB for BERT and ROBERTA models, and 5.2 GB for LLaMa-2 7B models. The decomposition part can be the most memory-consuming, as it involves reshaping a 3-dimensional tensor into a matrix with a dimension size equal to the product of two dimension sizes of the tensor.

However, this part is executed only once during the entire fine-tuning process and can be computed on the CPU in a short time. The Higher Order SVD decomposition of such objects takes approximately 78 seconds for BERT and RoBERTa_{base} layers and about 668 seconds for LLaMa on an Intel Xeon Gold 6342 CPU processor.

9 Ethics Statement

Our proposed approach involves a novel method for fine-tuning large language models, which can be considered as cost-effective as we only update 0.1% of the weights. This type of fine-tuning is environmentally friendly as it reduces resource wastage. We utilized pre-trained models from the Hugging Face repository and implemented updates using the Pytorch library. We exclusively used open-source datasets to avoid any potential harm or ethical concerns. By prioritizing ethical standards and recognizing potential risks, we strive to promote responsible and sustainable research practices.

A Appendix A

B Appendix B

C Appendix C

The average GLUE results for the BERT and RoBERTa_{base} models with respect to the number of remaining updated parameters in Linear layers. Tables 11 and 12 show that under the 0.8% of the remaining

²https://github.com/sayankotor/sparse_grads

Table 7: The training step execution speed, measured in steps per second (where a higher value indicates faster execution), is reported for the RoBERTa_{base} model. The last two rows describe the SparseGrad-Method with Sparse-by-Dense multiplication and with Regular matrix multiplication.

Method / Dataset	AVG	STSB	CoLA	MNLI	MRPC	QNLI	QQP	RTE	SST2
Regular FT	4.11	2.9	4.3	4.2	4.1	3.1	4.7	4.2	5.1
LoRA	4.7	2.8	5.8	6.2	6.3	3.4	4.1	3.2	4.4
SparseGrad, Sparse-by-Dense	4.3	3.8	1.8	3.9	3.1	3.5	5.6	6.3	6.2
SparseGrad, Regular	0.9	0.4	0.3	0.4	1.9	0.8	0.7	1.6	1.1

Table 8: Peak memory measurement in MB for training loop for the model RoBERTa_{base}.

Method / Dataset	AVG	STSB	CoLA	MNLI	MRPC	QNLI	QQP	RTE	SST2
Regular FT	1345	1344	1358	1350	1362	1369	1333	1314	1339
LoRA	944	969	978	986	998	938	935	902	855
SparseGrad, Sparse-by-Dense	1016	997	1082	1017	1110	1019	981	960	980
SparseGrad, Regular	1210	1283	1212	1256	1183	1245	1172	1116	1209

parameters, performance tends to decrease.

D Appendix D

Best training parameters for all models. In all experiments, we repeat fine-tuning 3 times over different seeds and report the average score.

E Appendix E

Responses from the models to an example from Inflection-Benchmarks are shown. While all models perform fairly well, the LoRA-trained model overlooks the fact that public speaking can be stressful for an introvert when answering the first question.

[Question 1]

Write a persuasive email to convince your introverted friend, who dislikes public speaking, to volunteer as a guest speaker at a local event. Use compelling arguments and address potential objections. Please be concise.

LLaMa2-7B Regular FT:

Here’s a persuasive email to convince your introverted friend to volunteer as a guest speaker at a local event: Subject: Join us as a guest speaker at our local event Dear (Friend’s Name), I hope this email finds you well. I hope you’re doing well! I’m writing to invite you to be a guest speaker at our upcoming local event. As a friend, I know how passionate you are!

LLaMa2-7B LoRA:

Subject Line: Please, Volunteer as a Guest Speaker! Dear John, I hope you are doing well. I am writing to you because I know you are an introverted person and do not like public speaking. However, I have a great opportunity for you to volunteer as a guest speaker at the event!

LLaMa2-7B SparseGrad:

Subject: Volunteer as a Guest Speaker at a Local Event Dear [Name], I hope this email finds you well. I wanted to reach out to you about an exciting opportunity to volunteer as a guest speaker at a local event. As you know, public speaking can be a nerve-wracking experience, but it’s also a great way to share your knowledge and expertise with others.

Table 9: Comparative results of BERT model for 20-epoch task-specific fine-tuning.

Method RTE	#Trainable SST2	Parameters	Model	MLP Layer	AVG	STSB	CoLA	MNLI	MR
Regular FT	109 mln	3 mln	82.5	89.3±.6	59.0±1.9	84.0±.3	86.2±1.1	89.3±1.3	91.1
LoRA	53 mln	0.03 mln	81.6	89.2±.7	58.4±2.3	84.2±.2	83.8±.6	89.3±.8	91.0
SparseGrad	53 mln	0.03 mln	82.6	89.2±.4	58.8±0	84.0±1.3	86.6±.5	89.4±1.6	90.9
MeProp	53 mln	0.03 mln	82.1	88.9±.5	58.4±.8	83.3±.3	84.2±.6	89.6±.3	90.4

Table 10: Comparative results of RoBERTa for 20-epoch task-specific fine-tuning.

Method RTE	#Trainable SST2	parameters	Model	MLP Layer	AVG	STSB	CoLA	MNLI	MRPC
Regular FT	125 mln.	3 mln.	84.2	90.4±.3	59.7±1.4	87.7±.1	90.0±.6	90.6±.8	91.5±.
LoRA	68 mln.	0.03 mln.	83.1	90.5±.2	60.6±1.7	87.5±.1	88.4±.6	90.0±.8	91.4±.
SparseGrad	68 mln.	0.03 mln.	83.6	90.8±.2	60.0±1.6	87.5±.1	89.6±1.1	91.5±.6	91.5±.
MeProp	68 mln.	0.03 mln.	82.5	90.7±.1	59.2±1.3	85.9±.1	89.1±0.9	89.4±.5	90.5±.

[Question 2]

What are some business etiquette norms when doing business in Japan?

LLaMa2-7B Regular FT:

If you're doing business in Japan, it's important to follow their business etiquette norms. Here are some key things to keep in mind: 1. Hierarchy: In Japan, it's important to respect the hierarchy in the workplace. Address your Japanese colleagues with the appropriate title, and don't speak to them unless they initiate the conversation. 2. Punctuality: Being punctual is considered very important in Japanese culture, so make sure to arrive on time for meetings and events. 3. Business Cards: Business cards are an important part of doing business in Japan. Make sure to have your own business card made beforehand and carry it with you at all times.

LLaMa2-7B LoRA:

Of course! Here are some business etiquette norms when doing business in Japan: 1. Business Cards: Business cards are a vital part of business interactions in Japan. Always have a few business cards with you and exchange them with others during introductions. 2. Punctuality: Time is sacred in Japan, and being punctual is essential for business meetings. Arrive on time or even a few minutes early. 3. Gifts: Gifts are common in business interactions in Japan, especially for special occasions. Small personalized gifts such as chocolates or flowers are appropriate.

LLaMa2-7B SparseGrad:

Here are some business etiquette norms to keep in mind when doing business in Japan: 1. Respect for seniority: In Japanese culture, age and seniority are highly respected, so it is important to show respect to those who are older or in a higher position than you. 2. Business cards: Business cards are an important part of doing business in Japanese culture, and it is customary to exchange business cards at the beginning of a meeting. 3. Gifts: Gifts are often given as a sign of appreciation or gratitude in Japanese culture. It is important to give gifts that are appropriate for the occasion and not too extravagant.

Table 11: GLUE score as a function of the weight gradient sparsity in BERT

Method	% of remained params in Linear Layers	AVG	STSB	CoLA	MNLI	MRPC	Q
SST2							
SparseGrad	100	82.6	89.2 $\pm$.4	58.8 $\pm$ 0	84.0 $\pm$ 1.3	86.6 $\pm$.5	89.2 $\pm$.1
92.4 $\pm$.1							
SparseGrad18k	0.8	81.5	89.1 $\pm$.3	59.1 $\pm$.5	83.8 $\pm$.1	84.6 $\pm$.8	89.2 $\pm$.1
92.4 $\pm$.6							
SparseGrad22k	1	82.2	89.7 $\pm$.6	60.0 $\pm$.4	83.9 $\pm$.1	84.6 $\pm$ 1.5	88.8 $\pm$.1
92.3 $\pm$.3							
SparseGrad30k	1.2	82.0	89.2 $\pm$.4	59.1 $\pm$.5	84.1 $\pm$.3	85.4 $\pm$.6	89.2 $\pm$.1
92.2 $\pm$.4							
SparseGrad100k	4.2	82.2	89.3 $\pm$.3	60.0 $\pm$.3	83.8 $\pm$.2	85.1 $\pm$ 1.2	88.8 $\pm$.1
92.4 $\pm$.3							

Table 12: GLUE score as a function of the weight gradient sparsity in RoBERTa

Method	% of remained params in Linear Layers	AVG	STSB	CoLA	MNLI	MRPC	Q
SST2							
SparseGrad	100	83.6	90.8 $\pm$.2	60.0 $\pm$ 1.6	87.5 $\pm$.1	89.6 $\pm$ 1.1	91.2 $\pm$.1
94.2 $\pm$.1							
SparseGrad18k	0.8	83.4	90.9 $\pm$.2	59.7 $\pm$.1	87.4 $\pm$.4	89.2 $\pm$.7	89.2 $\pm$.1
94.0 $\pm$.4							
SparseGrad22k	1	83.6	90.6 $\pm$.2	58.8 $\pm$.4	87.7 $\pm$.1	90.0 $\pm$.3	90.2 $\pm$.1
94.6 $\pm$.2							
SparseGrad30k	1.2	83.6	90.8 $\pm$.3	59.4 $\pm$.4	87.6 $\pm$.1	89.8 $\pm$.4	91.2 $\pm$.1
94.2 $\pm$.1							
SparseGrad100k	1.4	83.9	90.9 $\pm$.3	59.8 $\pm$.5	87.0 $\pm$.1	89.7 $\pm$.5	89.2 $\pm$.1
94.1 $\pm$.3							

Table 13: Best training parameters on GLUE benchmark for BERT model.

Dataset	batch size / learning rate
STSB	32 / 1.24e-4
CoLA	32 / 3.15e-5
MNLI	32 / 6.07e-6
MRPC	32 / 1.22e-5
QNLI	16 / 1.94e-5
QQP	32 / 1.41e-5
RTE	16 / 6.81e-5
SST2	32 / 1.47e-5

figures/original_linear.pdf

(a) Original Linear Layer

figures/sparsegrad_linear.pdf



Figure 2: Gradients on the 5-th BERT MLP: $U \frac{\partial L}{\partial W^T} V^T$ (right) is more sparse than the original $\frac{\partial L}{\partial W^T}$ (left).




Figure 3: Strided structure of $\frac{\partial L}{\partial Y}$ (left) and visualizations of % nonzero elements in $\frac{\partial L}{\partial Y}$ throughout training (right).

Table 14: Best training parameters on GLUE benchmark for RoBERTa model.

Dataset	batch size / learning rate
STSB	16 / 2.70e-5
CoLA	16 / 1.01e-5
MNLI	32 / 1.51e-5
MRPC	32 / 1.9e-5
QNLI	16 / 1.91e-5
QQP	16 / 5.11e-6
RTE	32 / 3.05e-5
SST2	16 / 1.33e-5

Table 15: Best training parameters on GLUE benchmark for RoBERTa-large model.

Dataset	batch size / learning rate
STSB	32 / 7.71e-5
CoLA	16 / 1.8e-5
MNLI	16 / 1.15e-6
MRPC	32 / 2.47e-5
QNLI	16 / 8.83e-6
QQP	32 / 7.2e-6
RTE	32 / 1.02e-5
SST2	32 / 1.02e-5